

# Neighborhood Label-Barycenter LMB Fusion for Distributed Multi-Target Tracking Under Unreliable Communication

FIRST AUTHOR, Member, IEEE  
Institution, City, Country

SECOND AUTHOR, Member, IEEE  
Institution, City, Country

THIRD AUTHOR, Member, IEEE  
Institution, City, Country

**Abstract**— Unreliable peer-to-peer communication can turn distributed labeled multi-Bernoulli (LMB) tracking into a component-correspondence failure. Neighboring filters may assign different labels to one target, or reuse one label for states that have drifted apart. Scalar arithmetic-average (AA) and Kullback–Leibler average (KLA) weights decide how much probability mass to trust. They do not decide which active Bernoulli output components share a label/moment correspondence. We introduce a graph-local neighborhood label-barycenter operator for distributed active LMB outputs. Each sensor selects a median-cardinality medoid reference, assigns neighboring tracks to that reference, and replaces matched Gaussian posteriors with a first-two-moment barycenter. A reference-only ablation separates label-set canonicalization from posterior averaging. In 50 paired trials under tiered packet loss, the operator reduces network optimal sub-pattern assignment (OSPA) disagreement by 81.59%. Relative to a fixed target-wise spatial-KLA AA baseline, it also reduces local expected OSPA (E-OSPA) by 17.15% and root-mean-square error (RMSE) by 6.35%. By contrast, the reference-only ablation gives only 0.54% RMSE reduction. A held-out 50-trial replication preserves this full-versus-reference separation: 6.64% RMSE reduction for the full operator versus 0.82% for reference-only label copying. Fixed-parameter checks under harsher packet loss, ring topology, and partial field of view (FOV) preserve the same separation. These results support an output-level correspondence-and-projection view: scalar fusion weights remain necessary, but matched barycenters supply the label/moment contract those weights assume.

Manuscript received Month 00, 2026; revised Month 00, 2026; accepted Month 00, 2026. This work was supported in part by [Funding Agency] under Grant [Grant Number]. Corresponding author: FIRST AUTHOR.

First Author, Second Author, and Third Author are with [Institution], [City], [Country] (e-mail: [author@example.com]).

This study uses simulated multi-target tracking data and no human-subject data.

Code and reproducibility artifacts are available at [repository DOI/URL].

0018-9251 © 2026 IEEE

## I. INTRODUCTION

DISTRIBUTED multi-target tracking must estimate a time-varying target set while keeping local sensor estimates mutually usable under lossy communication. In networked surveillance and cooperative sensing systems, a local track is useful to a neighbor only when its identity and state are comparable with that neighbor’s track. Labeled random finite sets (RFSs) and the labeled multi-Bernoulli (LMB) filter represent target birth, persistence, death, and track identity [1], [2]. In a peer-to-peer network, however, labels also encode local filtering history. After asymmetric measurements and packet drops, two sensors can assign different labels to the same target, or reuse one label for states that have drifted apart. The label set therefore becomes a correspondence contract: if that contract differs across sensors, a mathematically valid fusion rule can still combine the wrong Bernoulli components.

Distributed RFS fusion is usually built around conservative geometric rules, such as generalized covariance intersection or the Kullback–Leibler average (KLA). It is also built around arithmetic-average (AA) rules that preserve weak components and support heterogeneous labeled and unlabeled filters [3]–[6]. These rules allocate probability mass across neighbors once a component correspondence has been supplied. They do not, by themselves, decide which Bernoulli component in one local LMB posterior corresponds to which component in another. Once unmatched components enter the existence and spatial fusion consumers, better scalar weights can attenuate a bad contribution but cannot create the missing correspondence. The output can then keep a plausible cardinality while averaging states from different targets.

We target this residual correspondence failure directly while keeping the scope narrower than density-pooling design. Once target-wise scalar weights are routed to the existence and spatial fusion consumers, the open question is whether a graph-local projection can construct the missing component map and average only matched posterior states. We test this question against a fixed target-wise AA-LMB baseline in which fusion weights are resolved at the object level before the existence and spatial consumers. This choice prevents the proposed gain from being attributed to scalar-weight routing. The remaining design question is a correspondence-contract construction problem, not another scalar-weight search. It asks how to choose a neighborhood reference label set, establish assignments to that reference, and replace posterior states after correspondence is established.

We propose a neighborhood label-barycenter operator for distributed LMB fusion. Each sensor chooses a median-cardinality medoid reference from its communication neighborhood and matches neighboring tracks to the reference labels using the Hungarian assignment method [7]. It then replaces each matched Gaussian group by its empirical first two moments. The operator is an output-space projection: repeating it propagates label-space agreement over the graph without reading a global label set or replacing the upstream Bernoulli existence update.

Contributions. The contributions form a single mechanism chain. First, the paper identifies a residual AA-LMB failure mode in which target-wise scalar weights are routed correctly but component labels still disagree. Second, it proposes a graph-local medoid-assignment moment-barycenter operator that canonicalizes the neighborhood label set before averaging matched posterior states. Third, it states structural conditions for cardinality robustness, stable assignment, reference-label invariance, and centralized moment-consensus limits. Fourth, it validates the mechanism through paired 50-trial evidence and a held-out 50-trial replication. The evidence also includes fixed-parameter harsh-loss, topology-ring, and partial field-of-view (FOV) N50 checks, plus a reference-only ablation that separates label copying from matched posterior barycentering.

## II. RELATED WORK

Labeled RFS filters provide a Bayesian finite-set formulation for preserving target identity while estimating multi-object states [1]. The LMB filter gives a tractable approximation in which each Bernoulli component carries an existence probability and a labeled single-target density [2]. Distributed labeled-RFS fusion has been studied through generalized covariance-intersection fusion, robust distributed labeled-RFS fusion, consensus LMB updates, and resilient peer-to-peer LMB fusion [4], [5], [8], [9]. Recent LMB-specific studies have also examined multiview fusion in networks without feedback and efficient label matching for distributed LMB tracking [10], [11]. These methods address how sensors should share uncertain multi-object information, but practical labeled fusion still depends on whether components with different local labels are comparable.

KLA and logarithmic opinion-pool fusion are widely used because they avoid double counting under unknown correlations and connect naturally to consensus on probability densities [3]. This conservatism is valuable in distributed tracking, but geometric fusion can suppress weak components when local support is asymmetric. Arithmetic-average fusion follows a linear-pooling principle and has been analyzed statistically and information-theoretically for RFS densities [12]. Subsequent AA density-fusion work

extends this analysis to unlabeled and labeled RFS fusion and to heterogeneous unlabeled/labeled fusion. It also covers variational heterogeneous distributed fusion and weighted peer-to-peer RFS fusion [6], [13]–[15].

Heterogeneity and field-of-view mismatch create a separate difficulty. RFS fusion across heterogeneous sensors can be formulated without assuming identical local multi-object densities [16]. Labeled RFS fusion with different fields of view has also been studied explicitly [17]. These lines motivate robust fusion under nonidentical local views. They do not by themselves remove the output-level failure in which two local LMB filters carry different labels for the same physical target after packet loss.

Our method is complementary to AA, KLA, heterogeneous-fusion, and label-matching LMB work. The gap is not how to reweight an already matched density. It is how to construct the correspondence map that makes local Bernoulli components comparable before output fusion. The contribution is therefore not another density-pooling rule. Instead, it is a targeted output-space projection, orthogonal to those weighting choices. It chooses a graph-local reference label set and computes matched moment barycenters. Relative to label-matching and multiview LMB work, the present study isolates the active-output correspondence layer. It uses a reference-only ablation to separate label copying from matched posterior barycentering.

This distinction also shapes the evaluation. Finite-set fusion can have pointwise and cardinality behavior that move in different directions [18]. We therefore report both network disagreement and truth-referenced tracking error using optimal sub-pattern assignment (OSPA)-style metrics [19]. The method relies on assignment, for which the Hungarian method remains the standard polynomial-time solver [7].

## III. PROBLEM FORMULATION

Consider a network of  $S$  sensors. At time  $k$ , sensor  $s$  outputs an LMB estimate

$$X_s(k) = \{(\ell_{s,i}, r_{s,i}, \mu_{s,i}, \Sigma_{s,i})\}_{i=1}^{n_s(k)}, \quad (1)$$

where  $\ell_{s,i}$  is a label,  $r_{s,i} \in [0, 1]$  is existence probability, and  $(\mu_{s,i}, \Sigma_{s,i})$  are the first two moments of a Gaussian state approximation. Let  $\mathcal{N}_s$  denote the sensor indices available to sensor  $s$  through the communication graph, including  $s$  itself.

The core difficulty is that the sets  $\{\ell_{s,i}\}$  need not agree across sensors. A local filter may keep a label after a missed packet while a neighboring filter deletes it, or two sensors may assign different labels to the same physical target. Direct AA or KLA fusion of unmatched Bernoulli components then mixes component identities and states. We seek an estimate-

level operator

$$\mathcal{B}_s : \{X_j(k) : j \in \mathcal{N}_s\} \mapsto Y_s(k) \quad (2)$$

that is local to the graph, preserves an interpretable reference label set, and reduces both pairwise output disagreement and truth-referenced tracking error.

The problem is not that scalar AA/KLA weights are unimportant. They remain the mechanism for allocating trust across incoming densities once the components being compared are known. The missing object is a correspondence map from local components to a common reference. The operator below estimates that map from the output geometry before applying a moment projection. In the reported implementation, this projection is applied to active output tracks after the upstream AA existence update and thresholding. It therefore rewrites the label set and Gaussian moments of the surviving tracks; it does not replace the Bernoulli existence consumer.

#### IV. NEIGHBORHOOD LABEL-BARYCENTER FUSION

##### A. Reference Selection

For each sensor  $s$  and time  $k$ , the candidate neighborhood is  $\mathcal{N}_s$ . Let  $n_j(k)$  be the number of tracks output by neighbor  $j$ . The reference cardinality is chosen around the sample median,

$$m_s(k) = \text{median}\{n_j(k) : j \in \mathcal{N}_s\}. \quad (3)$$

The candidate set is

$$\mathcal{C}_s(k) = \arg \min_{j \in \mathcal{N}_s} |n_j(k) - m_s(k)|. \quad (4)$$

Among these candidates, the reference is the medoid under an OSPA-like set distance  $d_c$  with cutoff  $c$ :

$$\rho_s(k) = \arg \min_{j \in \mathcal{C}_s(k)} \frac{1}{|\mathcal{N}_s|} \sum_{q \in \mathcal{N}_s} d_c(X_j(k), X_q(k)). \quad (5)$$

The labels of  $X_{\rho_s(k)}(k)$  form the local reference label set. In the implementation,  $d_c$  is computed from position states by Hungarian matching, squared matched distances, and a cutoff penalty  $c^2$  for unmatched tracks, normalized by the larger cardinality. The cutoff is the configured consensus cutoff when available and otherwise follows the OSPA cutoff used by the tracking evaluation.

##### B. Assignment-Based Label Canonicalization

Let the reference have states  $\{\mu_{\rho,a}\}_{a=1}^m$ . For each neighbor  $j$ , we construct a cost matrix

$$C_{ba} = \|p(\mu_{j,b}) - p(\mu_{\rho,a})\|_2, \quad (6)$$

where  $p(\cdot)$  extracts position coordinates. The Hungarian method returns the minimum-cost assignment between neighbor states and reference states [7]. This assignment defines the matched group  $\mathcal{G}_{s,a}(k)$  for each reference label  $a$ .

##### C. Matched Moment Barycenter

For a matched group  $\mathcal{G}_{s,a}(k) = \{(\mu_i, \Sigma_i)\}_{i=1}^M$ , the output state is

$$\bar{\mu}_{s,a} = \frac{1}{M} \sum_{i=1}^M \mu_i, \quad (7)$$

$$\bar{\Sigma}_{s,a} = \frac{1}{M} \sum_{i=1}^M [\Sigma_i + (\mu_i - \bar{\mu}_{s,a})(\mu_i - \bar{\mu}_{s,a})^T]. \quad (8)$$

The output  $Y_s(k)$  keeps the reference active labels and replaces their Gaussian moments by  $(\bar{\mu}_{s,a}, \bar{\Sigma}_{s,a})$ . The projection does not estimate new Bernoulli existence probabilities. It is applied after upstream AA existence fusion and thresholding, passes the surviving active-track existence scores through with the selected reference labels, and rewrites only the label and moment fields. Equivalently, with  $M_i = \Sigma_i + \mu_i \mu_i^T$  and  $q_i = [\mu_i^T, \text{vec}(M_i)^T]^T$ , the operator averages first-two-moment coordinates. This is a moment-matching projection, not a covariance-consistency guarantee. The implementation repeats this neighborhood operator for  $H = 3$  rounds in the reported experiments.

##### D. Graph Locality and Computational Cost

The operator is graph-local at every round: sensor  $s$  only reads estimates indexed by  $\mathcal{N}_s$ , and no global label dictionary is constructed. The communicated object is the active-track tuple  $(\ell, r, \mu, \Sigma)$  already exchanged for output-level AA. The projection only adds local reference selection, assignment, and moment accumulation after reception. An  $H$ -round distributed implementation therefore requires  $H$  neighborhood output-exchange rounds and no centralized label service, while existence scores continue to come from upstream AA. If  $d_s = |\mathcal{N}_s|$  and  $n_{\max}$  upper bounds the number of tracks in the neighborhood, the dominant cost is assignment. The medoid stage evaluates set distances among neighborhood estimates, while the matching stage solves at most  $d_s$  rectangular assignments to the chosen reference. With a cubic Hungarian implementation, the per-round cost is bounded by  $O(d_s^2 n_{\max}^3 + d_s n_{\max}^3)$  in the worst case, plus linear moment accumulation over matched tracks. This cost is acceptable in the reported formation setting, but it explains why the method is slower than scalar-weight AA and motivates the runtime study reported with the experimental results.

#### V. STRUCTURAL PROPERTIES

The propositions below define the projection layer's algebraic contract. They separate guarantees of reference selection, assignment, and first-two-moment averaging from phenomena left to the underlying LMB/AA pipeline. All statements use fixed output estimates at one time step, after the upstream AA existence consumer has selected active tracks.




Fig. 1. Active-output correspondence projection. The input is the active neighborhood LMB output, not a global label dictionary. Fusion weights specify neighbor probability mass, but they do not define component correspondence across local LMB posteriors. The operator chooses a graph-local reference, matches neighbor tracks to that reference, passes upstream active-track existence scores through, and rewrites only labels and Gaussian moments.

| Validation-time projection contract |                                                              |
|-------------------------------------|--------------------------------------------------------------|
| Input:                              | $X_j(k)$ for $j \in \mathcal{N}_s$ , cutoff $c$ , rounds $H$ |
| 1 Reference:                        | median-cardinality OSPA medoid $\rho_s(k)$                   |
| 2 Match:                            | Hungarian assignment to reference labels                     |
| 3 Project:                          | moment barycenters for matched states                        |
| 4 Existence:                        | pass through upstream $r$ ; rewrite labels/moments only      |
| 5 Iterate:                          | local overwrite; no global label dictionary is read          |

Fig. 2. Validation-time implementation outline. The operator is an estimate-level projection layer: it consumes only graph-neighborhood outputs, does not read a global label set, and does not re-estimate Bernoulli existence probabilities.

#### Assumption 1 (Analysis scope):

Unless a proposition states otherwise, the structural claims below analyze a fixed neighborhood at one time step. They treat reference selection and assignment as operating on active tracks. They hold births and deaths outside the projection layer, and do not assume that close target crossings or cardinality mismatches are resolved by the barycenter itself.

#### Proposition 1 (Weighting is not matching):

Consider two sensors and two one-dimensional targets with distinct means  $x_1 \neq x_2$ . Sensor 1 labels them  $(a, x_1)$  and  $(b, x_2)$ , while sensor 2 labels them  $(a, x_2)$  and  $(b, x_1)$ . A label-keyed scalar fusion consumer fuses equal labels with weights  $\alpha_a, \alpha_b \in [0, 1]$  and produces means

$$\hat{x}_a = \alpha_a x_1 + (1 - \alpha_a) x_2, \quad \hat{x}_b = \alpha_b x_2 + (1 - \alpha_b) x_1. \quad (9)$$

If both sensors have positive weight for a component, the fused mean lies between the two physical targets rather than selecting the missing permutation. Choosing an endpoint weight avoids mixing only by discarding one sensor for that component. Thus

scalar weights can suppress or emphasize an incoming component, but they do not by themselves infer the correspondence  $a \leftrightarrow b$  created by the label swap.

Proof:

The displayed equations are the scalar-weighted means used by a label-keyed consumer in this two-component example. For  $0 < \alpha_a < 1$ ,  $\hat{x}_a$  is a strict convex combination of  $x_1$  and  $x_2$ . The same holds for  $\hat{x}_b$  when  $0 < \alpha_b < 1$ . Recovering an endpoint for a component requires setting the corresponding weight to 0 or 1. That choice selects one sensor's label assignment rather than estimating the cross-sensor permutation. ■

#### Proposition 2 (Median-cardinality robustness):

If more than half of the sensors in  $\mathcal{N}_s$  output the same cardinality  $n^*$  at time  $k$ , then the reference selected by the median-cardinality stage also has cardinality  $n^*$ .

Proof:

With a strict majority at  $n^*$ , the sample median of neighborhood cardinalities is  $n^*$ . The minimum value of  $|n_j(k) - m_s(k)|$  is therefore zero, and it is attained only by sensors with  $n_j(k) = n^*$ . The subsequent medoid step is restricted to this candidate set. ■

#### Proposition 3 (Sufficient assignment stability):

Let the reference positions be  $\{z_a\}_{a=1}^m$  with separation

$$\Delta = \min_{a \neq b} \|z_a - z_b\|_2. \quad (10)$$

Suppose a neighbor has the same cardinality and contains one state  $y_a$  for each reference target such that  $\|y_a - z_a\|_2 \leq \epsilon$  after the physical correspondence is indexed by  $a$ . If  $\epsilon < \Delta/(m+1)$ , then the minimum-sum Euclidean assignment to the reference is the physical correspondence.

Proof:

The physical correspondence has assignment cost at most  $m\epsilon$ . Any nonphysical permutation assigns at least one state  $y_a$  to a different reference  $z_b$ ,  $b \neq a$ . For that pair,

$$\|y_a - z_b\|_2 \geq \|z_a - z_b\|_2 - \|y_a - z_a\|_2 \geq \Delta - \epsilon. \quad (11)$$

All other assignment costs are nonnegative, so the nonphysical assignment costs at least  $\Delta - \epsilon$ . The condition  $\epsilon < \Delta/(m+1)$  gives  $\Delta - \epsilon > m\epsilon$ , hence every nonphysical assignment has larger cost. ■

#### Proposition 4 (Moment-space projection):

For a matched group  $\mathcal{G} = \{(\mu_i, \Sigma_i)\}_{i=1}^M$ , define  $M_i = \Sigma_i + \mu_i \mu_i^T$  and  $q_i = [\mu_i^T, \text{vec}(M_i)^T]^T$ . The coordinate mean  $\bar{q} = \frac{1}{M} \sum_i q_i$  uniquely minimizes  $\sum_i \|q - q_i\|_2^2$ .

Recovering  $\bar{\mu}$  and  $\bar{M}$  from  $\bar{q}$  gives

$$\bar{\Sigma} = \bar{M} - \bar{\mu}\bar{\mu}^T = \frac{1}{M} \sum_{i=1}^M [\Sigma_i + (\mu_i - \bar{\mu})(\mu_i - \bar{\mu})^T], \quad (12)$$

which is positive semidefinite when the incoming covariances are positive semidefinite. The replacement is therefore a first-two-moment least-squares representative of the matched group and gives the first two moments of the equally weighted matched Gaussian mixture  $\frac{1}{M} \sum_i \mathcal{N}(\mu_i, \Sigma_i)$ .

Proof:

The Euclidean least-squares objective in  $q$  is minimized by the sample mean, giving  $\bar{\mu} = \frac{1}{M} \sum_i \mu_i$  and  $\bar{M} = \frac{1}{M} \sum_i M_i$ . Substituting  $M_i = \Sigma_i + \mu_i \mu_i^T$  and subtracting  $\bar{\mu}\bar{\mu}^T$  gives the displayed covariance expression. Each summand is positive semidefinite. The mean coordinate also satisfies  $\sum_i \|\mu_i - z\|_2^2 = \sum_i \|\mu_i - \bar{\mu}\|_2^2 + M\|z - \bar{\mu}\|_2^2$ . The update therefore removes within-group mean disagreement after correspondence is fixed. For any copied reference moment vector  $q_r$ , the same bias-variance identity gives  $\sum_i \|q_i - q_r\|_2^2 = \sum_i \|q_i - \bar{q}\|_2^2 + M\|q_r - \bar{q}\|_2^2$ . This gives the reference-copying inequality: barycentering is therefore never worse than reference copying in within-group moment error. Equality requires the copied reference to equal  $\bar{q}$ . ■

Proposition 5 (Stable-matching consensus limit):

For one matched reference label, suppose the correspondence is fixed and correct over a connected graph, and let

$$q_s^h = [(\mu_s^h)^T, \text{vec}(M_s^h)^T]^T, \quad M_s^h = \Sigma_s^h + \mu_s^h (\mu_s^h)^T. \quad (13)$$

If  $q_s^{h+1} = \sum_j W_{sj} q_j^h$  with a primitive doubly stochastic matrix  $W$  respecting the graph, then all  $q_s^h$  converge to  $S^{-1} \sum_j q_j^0$ . The limiting mean and second moment therefore equal the centralized equal-weight moment barycenter.

Proof:

Stack the  $q_s^h$  vectors. Since  $W^h \rightarrow S^{-1} \mathbf{1}\mathbf{1}^T$ , every node converges to the average initial moment vector. Recovering  $\Sigma$  from  $M - \mu\mu^T$  gives the covariance. ■

Proposition 6 (Upstream AA existence convexity):

If an AA existence update uses nonnegative weights that sum to one, then the fused existence probability remains in the convex hull of the incoming existence probabilities.

Proof:

The fused value is a convex combination of values in  $[0, 1]$ . ■

Proposition 7 (Reference-label invariance):

For every sensor  $s$ , the label set of  $Y_s(k)$  is a subset of the label set produced by one sensor in  $\mathcal{N}_s$ . Moreover, if every estimate in  $\mathcal{N}_s$  has the same labels and identical matched Gaussian moments, then  $\mathcal{B}_s$  returns that estimate unchanged.

Proof:

The operator defines output labels only from the selected reference  $X_{\rho_s(k)}(k)$ , so it never creates a label outside a neighborhood estimate. If all neighborhood estimates share the same labels and moments, every admissible reference has the same label set. Each assignment is the identity up to ties among identical states, and each matched group contains identical moments. The barycenter equations therefore return the same mean and covariance. ■

Proposition 8 (Cardinality-equivalent control):

With identical active inputs and the same median-cardinality reference rule, the reference-only and full moment-barycenter projections produce identical per-sensor output counts after every projection round. Consequently, metrics depending only on output cardinalities are identical for the two arms on the same paired trials.

Proof:

In one round, both arms choose from the same neighborhood cardinality multiset and output one component for each selected reference label. Copying versus barycentering changes only moments. Equal count fields are therefore preserved from round to round, so induction over the fixed number of rounds proves the claim. ■

These properties define a projection-layer contract, not a blanket tracking-optimality claim. The median-cardinality and assignment statements identify when the reference mechanism preserves a dominant cardinality and a separated physical correspondence. The assignment-stability condition is sufficient rather than necessary, and it excludes cardinality mismatch, births, deaths, and close target crossings. The consensus-limit result concerns matched moment coordinates under stable matching and stochastic graph weights, whereas the implementation uses three finite rounds and recomputes references and assignments. The existence convexity statement belongs to the upstream AA existence consumer, which the projection layer leaves intact.

The empirical claim is correspondingly narrow. The barycenter-versus-copying identity is an internal moment-space statement under fixed correspondence; if the assignment is wrong, a lower within-group moment error can still describe the wrong group. The cardinality-equivalent control explains why cardinality dispersion and cardinality error cannot by themselves distinguish label copying from moment barycentering. We therefore test whether reference-



TABLE I  
Mechanism-isolation protocol for the AA comparison.

| Arm                     | AA route | Ref. label | Moment avg. |
|-------------------------|----------|------------|-------------|
| Fixed spatial-KLA<br>AA | on       | off        | off         |
| Reference-only          | on       | on         | off         |
| Label-barycenter        | on       | on         | on          |

only canonicalization reduces label-space disagreement, and whether matched moment barycenters add localization gains once labels are aligned. The experimental section holds scalar AA routing fixed across the baseline, reference-only, and full-projection arms.

## VI. EXPERIMENTAL SETUP

The evaluation uses an eight-sensor formation-tracking scenario with tiered fixed packet-loss probabilities. Each 50-trial paired run uses the same seeds and the same per-trial packet-loss assignments across methods. The loss levels are  $\{0, 0.1, 0.2, 0.5\}$ , with one sensor at 0, four sensors at 0.1, one sensor at 0.2, and two sensors at 0.5 in each trial. The LMB filters use AA parallel update mode, Metropolis fusion weights, a maximum of three Gaussian-mixture components, 150 loopy-belief-propagation iterations, and an existence threshold of 0.18.

We compare three fixed arms chosen to isolate mechanism rather than to search for scenario-specific hyperparameters. The fixed baseline is the target-wise spatial-KLA AA implementation with link-quality weighting, existence confidence, and structure-aware spatial decoupling. It tests whether label-space errors remain after scalar fusion weights reach the correct consumers. Raw validation reports retain an internal implementation label for provenance. For manuscript consistency, all manuscript tables and comparisons use the fixed spatial-KLA AA baseline name and a fixed parameterization rather than per-scenario retuning. The proposed arm keeps that upstream existence pathway and applies neighborhood label-barycenter over communication neighborhoods with  $H = 3$  fixed iterations. The ablation uses the same reference selection and label canonicalization but copies medoid reference states without matched-state barycenters. All arms keep fixed parameters and receive the same measurements and packet-loss realization in each paired trial.

We separate network-consensus outcomes from truth-referenced tracking outcomes. Network disagreement uses pairwise output OSPA, localization disagreement, and cardinality dispersion across sensors. Truth-referenced tracking uses local expected OSPA (E-OSPA), root-mean-square error (RMSE), and cardinality error averaged across sensors. OSPA consistently combines localization and cardinality er-

rors for multi-object filters [19]. We report means, 95% confidence intervals, paired reductions, wins out of 50 paired trials, and sign-test  $p$  values.

The experiment suite is organized around mechanism claims rather than parameter search. The primary 50-trial run tests whether correspondence projection improves a fixed AA baseline once scalar weight routing is handled. The reference-only ablation tests whether a common neighborhood label set suffices, or whether matched posterior moments are needed. Held-out and boundary runs are fixed-parameter checks. They vary seeds, packet-loss severity, topology, sensing geometry, or the full-neighborhood ceiling only after method settings are fixed.

With base seed 1, the 50 trial seeds are 2 through 51. Each trial samples one eight-sensor packet-loss vector from the tiered profile and reuses it for all arms. Intervals are mean  $\pm 1.96$  standard errors over trial-level means. Paired reductions use baseline-minus-candidate deltas, and wins and two-sided sign-test values are computed on nonzero paired deltas.

The operator parameters are fixed before the reported N50 evaluation. In particular, the neighborhood version uses  $H = 3$  graph-local rounds and the same existence threshold as the fixed AA baseline. There is no per-scenario search over projection cut-offs, barycenter weights, thresholds, or trial-specific label rules. Held-out base-seed runs are reserved for robustness checks and are not used to choose these settings.

Runtime is measured as wall-clock filter time from one GNU Octave 11.1.0 execution on an Apple M4 workstation with 16 GB memory. Absolute seconds characterize this MATLAB-compatible prototype. The relative column gives the intended paired cost comparison.

The numerical tables and Fig. 3 are generated from tracked reports rather than edited by hand. The reproducibility package records source paths, SHA-256 digests, machine-readable summaries, and a ledger separating paired AA, held-out, harsh-loss, topology/FOV/full-topology, and contextual GA evidence. A verifier recomputes network disagreement, runtime, and local metrics before the PDF is built.

## VII. RESULTS

The results follow one evidence chain. Table II tests network and local gains. Table III gives matched-seed GA context. Table IV and Fig. 3 separate label-only from matched-barycenter effects. Table V repeats the held-out mechanism test. The ordering keeps the primary paired AA comparison, contextual GA placement, mechanism ablation, and held-out replication distinct.

In the primary paired run, Table II shows lower communication-facing and truth-facing error

TABLE II

Primary 50-trial paired validation under tiered packet loss. Lower values are better; bold entries mark column-wise best values including ties. Values are trial means. The full operator improves both network agreement and local tracking quality, while the reference-only ablation isolates label-set canonicalization.

| Method                        | Net OSPA | Loc. disag. | Card. disp. | E-OSPA | RMSE  | CardErr |
|-------------------------------|----------|-------------|-------------|--------|-------|---------|
| Fixed spatial-KLA AA          | 1.683    | 1.473       | 0.035       | 2.030  | 3.683 | 0.088   |
| Neighborhood reference-only   | 1.030    | 0.892       | 0.021       | 1.911  | 3.663 | 0.077   |
| Neighborhood label-barycenter | 0.310    | 0.216       | 0.021       | 1.681  | 3.449 | 0.077   |

TABLE III

Matched-seed GA context for the same 50 trial seeds and tiered packet-loss profile. Lower values are better; bold entries mark column-wise best values. The GA rows are generated by the tracked GA validation path and are not used in the paired AA sign tests.

| Method                        | Net OSPA | Loc. disag. | Card. disp. | E-OSPA | RMSE  | CardErr |
|-------------------------------|----------|-------------|-------------|--------|-------|---------|
| GA +structure-aware KLA       | 1.761    | 1.463       | 0.084       | 2.213  | 4.228 | 0.204   |
| GA +FID-FIA refinement        | 1.848    | 1.958       | 0.081       | 2.180  | 4.695 | 0.151   |
| Neighborhood label-barycenter | 0.310    | 0.216       | 0.021       | 1.681  | 3.449 | 0.077   |

TABLE IV

Paired reductions relative to the fixed spatial-KLA AA baseline. Bracketed values are 95% confidence intervals for absolute paired reductions;  $p$  is the two-sided sign-test value on nonzero paired deltas.

| Metric            | Full reduction        | Full wins | Full $p$    | Ref.-only reduction   | Ref.-only wins | Ref. $p$    |
|-------------------|-----------------------|-----------|-------------|-----------------------|----------------|-------------|
| Network OSPA      | 81.59% [1.359, 1.386] | 50/50     | $< 10^{-3}$ | 38.78% [0.639, 0.666] | 50/50          | $< 10^{-3}$ |
| Loc. disagreement | 85.31% [1.238, 1.275] | 50/50     | $< 10^{-3}$ | 39.46% [0.566, 0.596] | 50/50          | $< 10^{-3}$ |
| Card. dispersion  | 40.03% [0.012, 0.016] | 50/50     | $< 10^{-3}$ | 40.03% [0.012, 0.016] | 50/50          | $< 10^{-3}$ |
| E-OSPA            | 17.15% [0.342, 0.355] | 50/50     | $< 10^{-3}$ | 5.83% [0.112, 0.125]  | 50/50          | $< 10^{-3}$ |
| RMSE              | 6.35% [0.201, 0.267]  | 48/50     | $< 10^{-3}$ | 0.54% [-0.017, 0.056] | 27/50          | 0.672       |
| Card. error       | 12.27% [0.009, 0.013] | 46/50     | $< 10^{-3}$ | 12.27% [0.009, 0.013] | 46/50          | $< 10^{-3}$ |

for the full operator. For network OSPA, the 95% confidence intervals are [1.667, 1.699] for the fixed baseline, [0.300, 0.319] for the full method, and [1.014, 1.046] for the reference-only ablation. For local E-OSPA, the corresponding intervals are [2.013, 2.047], [1.664, 1.699], and [1.893, 1.929]. The separation is visible in both output agreement and truth-referenced tracking before paired tests are considered.

Table III places the AA projection beside two GA reference rows generated with the same seeds and tiered packet-loss profile. Because those rows come from the GA validation path, they are contextual baselines rather than paired significance-test inputs. The neighborhood label-barycenter row is lower than both GA rows on all six metrics, but no AA-versus-GA significance claim is inferred. The mechanism claim remains tied to the paired AA ablation.

Table IV and Fig. 3 separate label-set effects from state-barycenter effects. The reference-only ablation reduces network disagreement and shares the cardinality reductions of the full operator, as expected from the shared reference-cardinality projection. Its RMSE reduction is only 0.54%, with a confidence interval crossing zero, 27 wins out of 50 paired trials, and

no sign-test support. The full operator reaches 6.35% RMSE reduction, with 48 wins out of 50 and  $p < 10^{-3}$ . The full method also improves local E-OSPA in all 50 paired trials. Thus the paired rows attribute statistically supported localization gains to matched posterior barycenters, not to label canonicalization alone.

The held-out base-seed run repeats the mechanism test without changing the method parameters. On this run, the full operator reduces network OSPA by 81.85%, local E-OSPA by 17.15%, and local RMSE by 6.64% relative to the fixed spatial-KLA AA baseline. The full RMSE reduction has 48/50 wins with sign-test  $p < 10^{-3}$ , whereas the reference-only RMSE reduction is 0.82% with 28/50 wins and sign-test  $p = 0.480$ . The held-out evidence therefore preserves the same separation between label-set copying and matched posterior barycentering.

The projection has a measurable runtime cost. Full and reference-only arms cost 1.644 and 1.534 times the fixed baseline, respectively. The full method increases runtime from 42.487 s to 69.729 s. Assignment and output rewriting, rather than moment computation alone, dominate this cost.




Fig. 3. Mechanism separation in paired percentage reductions. Values are the paired reductions in Table IV relative to the fixed spatial-KLA AA baseline. Blue bars show the full label-barycenter operator, and gray-blue bars show the reference-only ablation. The cardinality rows are shared reference-cardinality controls, not moment-barycenter effects. Spatial E-OSPA/RMSE separation appears only after matched posterior barycenters are added.

TABLE V

Held-out N50 replication under the same fixed three-arm protocol. Reductions are relative to the fixed spatial-KLA AA baseline; bracketed intervals are 95% confidence intervals for absolute paired reductions. The run uses base seed 11.

| Metric | Full red. [CI]        | Full wins/ <i>p</i>       | Ref. red. [CI]        | Ref. wins/ <i>p</i>       |
|--------|-----------------------|---------------------------|-----------------------|---------------------------|
| OSPA   | 81.85% [1.367, 1.394] | 50/50, < 10 <sup>-3</sup> | 39.15% [0.647, 0.674] | 50/50, < 10 <sup>-3</sup> |
| E-OSPA | 17.15% [0.342, 0.357] | 50/50, < 10 <sup>-3</sup> | 5.88% [0.114, 0.126]  | 50/50, < 10 <sup>-3</sup> |
| RMSE   | 6.64% [0.210, 0.281]  | 48/50, < 10 <sup>-3</sup> | 0.82% [-0.008, 0.069] | 28/50, 0.480              |

## VIII. DISCUSSION AND LIMITATIONS

Taken together, the paired and boundary evidence support a specific interpretation rather than a generic fusion-weight claim. When target-wise AA weight routing is held fixed, reference-only projection reduces network disagreement and captures the shared cardinality effect. It does not produce a statistically supported RMSE separation in the primary or held-out mechanism tests. Adding matched moment barycenters creates the spatial tracking separation across the primary, held-out, harsh-loss, topology-ring, and partial-FOV checks. The full-topology run is an equivalence boundary, not a gain claim. The method is therefore a correspondence-and-projection layer complementary to AA/KLA weighting, not a replacement for density pooling.

This interpretation has boundaries. The operator remains an estimate-level projection layer. It is graph-local because each sensor uses only its communication neighborhood, but it is not a full recursive message update inside the LMB filtering loop. A recursive implementation would also need lifecycle guards for birth, death, and intermittently matched labels, because unstable matches can turn a useful barycenter into an erroneous track merge.

The main failure mode is assignment ambiguity rather than packet loss alone. When targets are closely spaced, cardinalities lack a neighborhood majority, or one local posterior is biased, the medoid can select a plausible but wrong correspondence. A recursive deployment should therefore gate projection by assignment-margin, covariance, track-age, or reli-

ability evidence, and fall back to reference-only or upstream AA output when correspondence is not credible. We omit such guards here to isolate the label-map and moment-barycenter mechanisms.

The method uses Euclidean position matching and equally weighted moment barycenters. These choices keep the operator interpretable and make the reference-only ablation clean, but they do not use covariance quality, label age, or link reliability inside the barycenter itself. A covariance-aware or reliability-weighted variant is a natural extension when local uncertainty differs strongly across sensors. The main validation is concentrated on a tiered packet-loss formation scenario.

The harsh packet-loss check kept the same fixed parameters and evaluated  $p_{\text{drop}} \in [0.2, 0.35, 0.5, 0.7]$  with sensor counts [1, 3, 2, 2]. Under this pre-specified setting, the full operator preserves the full-versus-reference separation: network OSPA falls by 79.57%, local E-OSPA by 17.06%, and RMSE by 7.98%, whereas reference-only RMSE falls by 2.47%.

Fixed-parameter N50 scenario-family checks probe sparse topology, partial-field-of-view sensing, and full-neighborhood ceiling. Full-versus-reference RMSE separation persists in topology-ring (RMSE 13.12% vs reference-only 5.17%); partial-FOV (RMSE 6.50% vs reference-only 2.13%). The full-topology check gives zero network disagreement and identical local E-OSPA/RMSE/CardErr (1.634/3.277/0.073), marking an equivalence boundary rather than a gain claim.

These checks broaden packet-loss severity, sparse-topology coverage, and partial-field-of-view sensing. They still preserve formation-family assumptions and



should not be read as a substitute for maneuvering-target or covariance-consistency studies.

## IX. CONCLUSION

This paper isolates a residual AA-LMB fusion error: scalar weights can reach the intended consumers while active LMB components still lack a shared cross-sensor correspondence. The proposed neighborhood label-barycenter operator chooses a median-cardinality medoid reference, assigns neighboring tracks to that reference, and computes first-two-moment barycenters for matched states. The paired primary and held-out N50 runs show large reductions in network disagreement and a repeatable RMSE separation between the full operator and reference-only copying. These results support the central mechanism: label canonicalization stabilizes the output label set, while matched posterior barycentering supplies the spatial gain. The current contribution is an output-level correspondence module; recursive LMB use still requires lifecycle and consistency guards.

## ACKNOWLEDGMENT

Supported in part by [Funding Agency] under Grant [Grant Number]. OpenAI Codex assisted with code navigation, log organization, evidence checks, and language polishing; the authors remain responsible for all technical content.

## REFERENCES

- [1] B.-N. Vo, B.-T. Vo, and D. Phung  
Labeled random finite sets and the Bayes multi-target tracking filter  
IEEE Transactions on Signal Processing, vol. 62, no. 24, pp. 6554–6567, 2014.
- [2] S. Reuter, B.-T. Vo, B.-N. Vo, and K. Dietmayer  
The labeled multi-bernoulli filter  
IEEE Transactions on Signal Processing, vol. 62, no. 12, pp. 3246–3260, 2014.
- [3] G. Battistelli and L. Chisci  
Kullback-leibler average, consensus on probability densities, and distributed state estimation with guaranteed stability  
Automatica, vol. 50, no. 3, pp. 707–718, 2014.
- [4] B. Wang, W. Yi, R. Hoseinnezhad, S. Li, L. Kong, and X. Yang  
Distributed fusion with multi-bernoulli filter based on generalized covariance intersection  
IEEE Transactions on Signal Processing, vol. 65, no. 1, pp. 242–255, 2017.
- [5] S. Li, W. Yi, R. Hoseinnezhad, G. Battistelli, B. Wang, and L. Kong  
Robust distributed fusion with labeled random finite sets  
IEEE Transactions on Signal Processing, vol. 66, no. 2, pp. 278–293, 2018.
- [6] T. Li, R. Yan, K. Da, and H. Fan  
Arithmetic average density fusion–part III: Heterogeneous unlabeled and labeled RFS filter fusion  
IEEE Transactions on Aerospace and Electronic Systems, vol. 60, no. 1, pp. 1023–1034, 2024.
- [7] H. W. Kuhn  
The hungarian method for the assignment problem  
Naval Research Logistics Quarterly, vol. 2, no. 1–2, pp. 83–97, 1955.
- [8] K. Shen, P. Dong, Z. Jing, and H. Leung  
Consensus-based labeled multi-bernoulli filter for multitarget tracking in distributed sensor network  
IEEE Transactions on Cybernetics, vol. 52, no. 12, pp. 12 722–12 733, 2022.
- [9] L. Gao, G. Battistelli, and L. Chisci  
Resilient labeled multi-bernoulli fusion with peer-to-peer sensor network  
Information Fusion, vol. 100, p. 101965, 2023.
- [10] Y. Jin and J. Li  
Multiview fusion with the labeled multi-bernoulli densities in the network without feedback  
IEEE Transactions on Aerospace and Electronic Systems, vol. 59, no. 6, pp. 9650–9664, 2023.
- [11] C. Ding, C. Shao, S. Zhou, D. Zhou, R. Du, and J. Liu  
Distributed multi-target tracking with labeled multi-bernoulli filter considering efficient label matching  
Frontiers of Information Technology & Electronic Engineering, vol. 26, no. 3, pp. 400–414, 2025.
- [12] T. Li, Y. Song, E. Song, and H. Fan  
Arithmetic average density fusion–part I: Some statistic and information-theoretic results  
Information Fusion, vol. 104, p. 102199, 2024.
- [13] T. Li  
Arithmetic average density fusion–part II: Unified derivation for unlabeled and labeled RFS fusion  
IEEE Transactions on Aerospace and Electronic Systems, vol. 60, no. 3, pp. 3255–3268, 2024.
- [14] T. Li, H. Liang, G. Li, J. Garcia Herrero, and Q. Pan  
Arithmetic average density fusion–part IV: Distributed heterogeneous fusion of RFS and LRFS filters via variational approximation  
IEEE Transactions on Signal Processing, vol. 73, pp. 1454–1469, 2025.
- [15] L. Zheng, L. Shi, F. Yang, and Y. Cai  
Distributed weighted fusion of RFS densities for peer-to-peer sensor networks  
Information Fusion, vol. 129, p. 104060, 2026.
- [16] W. Yi and L. Chai  
Heterogeneous multi-sensor fusion with random finite set multi-object densities  
IEEE Transactions on Signal Processing, vol. 69, pp. 3399–3414, 2021.
- [17] L. Gao, G. Battistelli, and L. Chisci  
Fusion of labeled RFS densities with different fields of view  
IEEE Transactions on Aerospace and Electronic Systems, vol. 58, no. 6, pp. 5908–5924, 2022.
- [18] M. Üney, J. Houssineau, E. Delande, S. J. Julier, and D. E. Clark  
Fusion of finite-set distributions: Pointwise consistency and global cardinality  
IEEE Transactions on Aerospace and Electronic Systems, vol. 55, no. 6, pp. 2759–2773, 2019.
- [19] D. Schuhmacher, B.-T. Vo, and B.-N. Vo  
A consistent metric for performance evaluation of multi-object filters  
IEEE Transactions on Signal Processing, vol. 56, no. 8, pp. 3447–3457, 2008.