

Neighborhood Label-Barycenter LMB Fusion for Distributed Multi-Target Tracking Under Unreliable Communication

FIRST AUTHOR, Member, IEEE
Institution, City, Country

SECOND AUTHOR, Member, IEEE
Institution, City, Country

THIRD AUTHOR, Member, IEEE
Institution, City, Country

Abstract— Distributed labeled multi-Bernoulli filters are attractive for multi-sensor multi-target tracking because each sensor can maintain a labeled posterior while exchanging compact information with neighboring sensors. Under unreliable communication, however, local filters can keep different label sets for the same physical targets, and average fusion then mixes posterior components that are not label-aligned. This paper develops a neighborhood label-barycenter operator for distributed LMB fusion: each sensor chooses a median-cardinality medoid reference from its communication neighborhood, matches neighboring tracks to the reference by assignment, and replaces matched posteriors by a first-two-moment barycenter over the matched group. The construction separates label canonicalization from posterior averaging, which makes it possible to test whether performance gains come from label-set agreement alone or from matched spatial barycenters. In a 50-trial tiered packet-loss formation scenario, the proposed operator reduces network OSPA disagreement by 81.59%, local E-OSPA by 17.15%, and local RMSE by 6.35% relative to a tuned spatial-KLA AA baseline, while a reference-only ablation gives only 0.54% RMSE reduction.

Index Terms— Arithmetic average fusion, distributed sensor networks, labeled multi-Bernoulli filter, multi-target tracking, random finite sets.

Manuscript received Month 00, 2026; revised Month 00, 2026; accepted Month 00, 2026. This paragraph will be finalized by IEEE and will include funding information when available. Corresponding author: FIRST AUTHOR.

First Author, Second Author, and Third Author are with Institution, City, Country (e-mail: author@example.com). Author information is provisional for the draft version.

This manuscript uses only simulated multi-target tracking data. No human or animal subjects are involved. AI-assisted drafting and editing will be disclosed in the acknowledgments in the final submission.

Supplementary material, code, and reproducibility artifacts will be provided upon submission if permitted by the journal.

0018-9251 © 2026 IEEE

I. INTRODUCTION

DISTRIBUTED multi-target tracking has to solve two coupled problems at once: estimating an unknown and time-varying set of target states, and keeping the local estimates at different sensors mutually usable when communication is lossy. Labeled random finite sets (RFSs) and the labeled multi-Bernoulli filter provide a principled representation for target birth, persistence, death, and track identity [1], [2]. In a centralized or well-synchronized system, labels make posterior components interpretable as tracks. In a peer-to-peer sensor network, the same labels can become a liability: local filters may instantiate, delete, or maintain labels differently after asymmetric measurements and packet drops.

Density fusion for distributed RFS tracking is usually built around conservative geometric rules, such as generalized covariance intersection or the Kullback–Leibler average (KLA), or around arithmetic-average (AA) fusion rules that preserve weak components and support heterogeneous labeled and unlabeled filters [3]–[6]. These rules define how to combine densities, but their practical LMB implementations still need a correspondence between local Bernoulli components. When the correspondence is wrong, a refined fusion weight cannot repair the fact that two local filters are speaking different label languages.

This paper starts from that failure mode. We use a tuned AA-LMB baseline in which target-wise weights are consumed consistently by the existence and spatial branches, so the comparison is not confounded by a scalar weight-routing bug. The remaining failure blocks contain two structural errors: label-set splits across local filters and spatial dispersion among tracks that should represent the same target. We therefore treat the problem as one of label canonicalization followed by matched posterior barycentering, rather than as another scalar weight-search problem.

We propose a neighborhood label-barycenter operator for distributed LMB fusion. At each time step, each sensor uses only the estimates available in its communication neighborhood. The sensor first chooses a reference estimate by a median-cardinality medoid rule, which prevents minority cardinality outliers from directly selecting the reference. It then matches neighboring estimates to the reference labels using the Hungarian assignment method [7]. For each reference label, the matched Gaussian states are replaced by the mean and covariance of the empirical matched mixture. Repeating this local operator over the graph propagates label-space agreement without reading a global label set.

The contribution is threefold.

- We identify a concrete failure mechanism for AA-LMB fusion under unreliable communication: label-space disagreement and matched-state

dispersion remain even after target-wise weight handling is corrected.

- We introduce a graph-local label-barycenter operator that combines median-cardinality medoid reference selection, assignment-based label canonicalization, and matched first-two-moment posterior barycenters.
- We provide structural analysis and paired experiments showing that the barycenter step, not label-set copying alone, is responsible for the main spatial tracking gain.

II. RELATED WORK

Labeled RFS filters provide a Bayesian finite-set formulation for preserving target identity while estimating multi-object states [1]. The LMB filter gives a tractable approximation in which each Bernoulli component carries an existence probability and a labeled single-target density [2]. Multi-sensor and distributed versions of labeled RFS filters have been studied through generalized covariance-intersection fusion, robust distributed labeled-RFS fusion, consensus LMB updates, and resilient peer-to-peer LMB fusion [4], [5], [8], [9]. These methods address how sensors should share uncertain multi-object information, but practical labeled fusion still depends on whether components with different local labels are comparable.

KLA and logarithmic opinion-pool fusion are widely used because they avoid double counting under unknown correlations and connect naturally to consensus on probability densities [3]. This conservatism is valuable in distributed tracking, but geometric fusion can suppress weak components when local support is asymmetric. Arithmetic-average fusion follows a linear-pooling principle and has been analyzed statistically and information-theoretically for RFS densities [10]. Subsequent AA density-fusion work gives unified derivations for unlabeled and labeled RFS fusion, heterogeneous unlabeled/labeled fusion, and variational heterogeneous distributed fusion [6], [11], [12]. Fisher-information weighted AA fusion for multi-rate filters and weighted peer-to-peer RFS fusion further show that recent work is moving from one-shot density pooling toward communication- and filter-quality-aware AA designs [13], [14].

Heterogeneity and field-of-view mismatch create a separate difficulty. RFS fusion across heterogeneous sensors can be formulated without assuming identical local multi-object densities [15], and labeled RFS fusion with different fields of view has been studied explicitly [16]. Information-geometric labeled-RFS fusion offers another route for distributed heterogeneous state estimation [17]. These lines motivate robust fusion under nonidentical local views, but they do not by themselves remove the output-level failure in which two local LMB filters carry different labels for the same physical target after packet loss.

Our method is complementary to AA, KLA, and information-geometric weighting. It does not replace the underlying density-fusion rule. Instead, it supplies a graph-local label-aligned projection layer that makes local Bernoulli components comparable before the fused output is evaluated. Finite-set fusion also raises consistency questions that are not present in single-target Gaussian fusion: pointwise consistency and global cardinality can move in different directions [18]. We therefore report both network disagreement and truth-referenced tracking error using OSPA-style metrics [19]. The method relies on assignment, for which the Hungarian method remains the standard polynomial-time solver [7], and on graph-local iteration, which is closely related to average-consensus ideas in sensor networks [20].

III. PROBLEM FORMULATION

Consider a network of S sensors. At time k , sensor s outputs an LMB estimate

$$X_s(k) = \{(\ell_{s,i}, r_{s,i}, \mu_{s,i}, \Sigma_{s,i})\}_{i=1}^{n_s(k)}, \quad (1)$$

where $\ell_{s,i}$ is a label, $r_{s,i} \in [0, 1]$ is existence probability, and $(\mu_{s,i}, \Sigma_{s,i})$ are the first two moments of a Gaussian state approximation. Let $\mathcal{N}_s$ denote the sensor indices available to sensor s through the communication graph, including s itself.

The core difficulty is that the sets $\{\ell_{s,i}\}$ need not agree across sensors. A local filter may keep a label after a missed packet while a neighboring filter deletes it, or two sensors may assign different labels to the same physical target. Direct AA or KLA fusion of unmatched Bernoulli components then mixes component identities and states. We seek an estimate-level operator

$$\mathcal{B}_s : \{X_j(k) : j \in \mathcal{N}_s\} \mapsto Y_s(k) \quad (2)$$

that is local to the graph, preserves an interpretable reference label set, and reduces both pairwise output disagreement and truth-referenced tracking error.

IV. NEIGHBORHOOD LABEL-BARYCENTER FUSION

A. Reference Selection

For each sensor s and time k , the candidate neighborhood is $\mathcal{N}_s$. Let $n_j(k)$ be the number of tracks output by neighbor j . The reference cardinality is chosen around the sample median,

$$m_s(k) = \text{median}\{n_j(k) : j \in \mathcal{N}_s\}. \quad (3)$$

The candidate set is

$$\mathcal{C}_s(k) = \arg \min_{j \in \mathcal{N}_s} |n_j(k) - m_s(k)|. \quad (4)$$

Among these candidates, the reference is the medoid under an OSPA-like set distance d_c with cutoff c :

$$\rho_s(k) = \arg \min_{j \in \mathcal{C}_s(k)} \frac{1}{|\mathcal{N}_s|} \sum_{q \in \mathcal{N}_s} d_c(X_j(k), X_q(k)). \quad (5)$$

Inputs: local LMB estimates $\{X_j(k) : j \in \mathcal{N}_s\}$, cutoff c , rounds H .
 For each round $h = 1, \dots, H$ and sensor s :
 1) compute neighborhood cardinalities and select the median-cardinality candidate set $\mathcal{C}_s(k)$;
 2) choose the OSPA-medoid reference $\rho_s(k)$ inside $\mathcal{C}_s(k)$;
 3) solve a Hungarian assignment from each neighbor output to the reference output using position costs;
 4) for every reference label, form the matched group and replace its state by the first-two-moment barycenter;
 5) write the resulting labeled estimate as the next-round output of sensor s .
 Output: label-aligned estimate $Y_s(k)$ for every sensor.

Fig. 1. Neighborhood label-barycenter operator used in the validation. The operator is an estimate-level projection layer: it consumes only graph-neighborhood outputs and does not read a global label set.

The labels of $X_{\rho_s(k)}(k)$ form the local reference label set.

B. Assignment-Based Label Canonicalization

Let the reference have states $\{\mu_{\rho,a}\}_{a=1}^m$. For each neighbor j , we construct a cost matrix

$$C_{ba} = \|p(\mu_{j,b}) - p(\mu_{\rho,a})\|_2, \quad (6)$$

where $p(\cdot)$ extracts position coordinates. The Hungarian method returns the minimum-cost assignment between neighbor states and reference states [7]. This assignment defines the matched group $\mathcal{G}_{s,a}(k)$ for each reference label a .

C. Matched Moment Barycenter

For a matched group $\mathcal{G}_{s,a}(k) = \{(\mu_i, \Sigma_i)\}_{i=1}^M$, the output state is

$$\bar{\mu}_{s,a} = \frac{1}{M} \sum_{i=1}^M \mu_i, \quad (7)$$

$$\bar{\Sigma}_{s,a} = \frac{1}{M} \sum_{i=1}^M [\Sigma_i + (\mu_i - \bar{\mu}_{s,a})(\mu_i - \bar{\mu}_{s,a})^T]. \quad (8)$$

The output $Y_s(k)$ keeps the reference labels and replaces their Gaussian moments by $(\bar{\mu}_{s,a}, \bar{\Sigma}_{s,a})$. The implementation repeats this neighborhood operator for $H = 3$ rounds in the reported experiments.

V. STRUCTURAL PROPERTIES

Proposition 1 (Median-cardinality robustness):

If more than half of the sensors in $\mathcal{N}_s$ output the same cardinality n^* at time k , then the reference selected




Fig. 2. The proposed operator separates label canonicalization from state fusion. Each sensor chooses a robust local reference, assigns neighbor tracks to the reference labels, and computes matched posterior barycenters. Repeated neighborhood rounds propagate label-space agreement without global label access.

by the median-cardinality stage also has cardinality n^* .

Proof:

With a strict majority at n^* , the sample median of neighborhood cardinalities is n^* . The minimum value of $|n_j(k) - m_s(k)|$ is therefore zero, and it is attained only by sensors with $n_j(k) = n^*$. The subsequent medoid step is restricted to this candidate set. ■

Proposition 2 (Moment projection):

For a matched group $\mathcal{G} = \{(\mu_i, \Sigma_i)\}_{i=1}^M$, the barycenter mean minimizes $\sum_i \|z - \mu_i\|_2^2$, and $(\bar{\mu}, \bar{\Sigma})$ are the first two moments of the equally weighted Gaussian mixture $M^{-1} \sum_i \mathcal{N}(\mu_i, \Sigma_i)$.

Proof:

The least-squares objective has gradient $2Mz - 2 \sum_i \mu_i$, hence its unique minimizer is $M^{-1} \sum_i \mu_i$. The covariance identity follows from total covariance:

$$\text{Cov}(x) = \mathbb{E}_i[\text{Cov}(x|i)] + \text{Cov}_i[\mathbb{E}[x|i]]. \quad (9)$$

Proposition 3 (Existence convexity):

If an AA existence update uses nonnegative weights that sum to one, then the fused existence probability remains in the convex hull of the incoming existence probabilities.

Proof:

The fused value is a convex combination of values in $[0, 1]$. ■

Proposition 4 (Reference-label invariance):

For every sensor s , the label set of $Y_s(k)$ is a subset of the label set produced by one sensor in $\mathcal{N}_s$. Moreover, if every estimate in $\mathcal{N}_s$ has the same labels and identical matched Gaussian moments, then $\mathcal{B}_s$ returns that estimate unchanged.

Proof:

The operator defines output labels only from the selected reference $X_{\rho_s(k)}(k)$, so it never creates a label outside a neighborhood estimate. If all neighborhood estimates share the same labels and moments, every admissible reference has the same label set, each assignment is the identity up to ties among identical states, and each matched group contains identical moments. The barycenter equations therefore return the same mean and covariance. ■

These properties are structural. They do not imply that tracking error must decrease for every scene. The empirical section therefore tests the operator against a tuned spatial-KLA AA baseline and a reference-only ablation.

VI. EXPERIMENTAL SETUP

The evaluation uses an eight-sensor formation-tracking scenario with tiered fixed packet-loss probabilities. Each 50-trial paired run uses the same seeds and the same per-trial packet-loss assignments across methods. The loss levels are $\{0, 0.1, 0.2, 0.5\}$, with one sensor at 0, four sensors at 0.1, one sensor at 0.2, and two sensors at 0.5 in each trial. The LMB filters use AA parallel update mode, Metropolis fusion weights, maximum three Gaussian-mixture components, 150 loopy-belief-propagation iterations, and an existence threshold of 0.18.

We compare three arms. The baseline is a tuned spatial-KLA AA implementation with target-wise weight handling, link-quality weighting, existence confidence, and structure-aware spatial decoupling. The proposed arm enables neighborhood label-barycenter projection over communication neighborhoods with $H = 3$ iterations. The ablation arm uses the same reference selection and label canonicalization but copies the medoid reference states without matched-state barycenters.

Network disagreement is measured by pairwise output OSPA, localization disagreement, and cardinality dispersion across sensors. Truth-referenced tracking quality is measured by local E-OSPA, RMSE, and cardinality error averaged across sensors. OSPA is used because it consistently combines localization and cardinality errors for multi-object filters [19]. We report mean values, 95% confidence intervals over trials, paired reductions relative to the tuned baseline, wins out of 50 paired trials, and sign-test p values.

VII. RESULTS

Table I shows that the proposed operator improves the communication-facing and truth-facing metrics simultaneously. The 95% confidence intervals for network OSPA are $[1.666520, 1.698693]$ for the tuned baseline, $[0.300179, 0.319457]$ for the full method, and $[1.014315, 1.045808]$ for the reference-




Fig. 3. Paired percentage reductions over the tuned spatial-KLA AA baseline. The reference-only ablation captures the effect of choosing and copying a consistent local reference label set. The full method gives substantially larger localization and RMSE gains because it also averages matched posterior states.

only ablation. For local E-OSPA, the corresponding intervals are $[2.012605, 2.046677]$, $[1.663594, 1.699373]$, and $[1.893427, 1.929145]$.

Table II and Fig. 3 separate label-set effects from state-barycenter effects. The reference-only ablation reduces network disagreement because sensors in a neighborhood move toward a common label set, but its RMSE reduction is only 0.54%, with a confidence interval crossing zero and 27 wins out of 50 paired trials. The full operator reaches 6.35% RMSE reduction, with 48 wins out of 50 and a sign-test value of $p = 2.267 \times 10^{-12}$. The full method also improves local E-OSPA in all 50 paired trials and cardinality error in 46 out of 50 paired trials. These results support the claim that matched posterior barycenters, not label copying alone, drive the spatial tracking gain.

The computational cost is higher than the tuned baseline. Mean filter runtime increases from 62.245 s to 102.347 s, or 1.647 times the tuned baseline, for the 50-trial validation (Table III). The reference-only ablation costs 1.542 times the baseline, so most of the overhead comes from repeated neighborhood matching and output rewriting rather than from moment computation alone.

VIII. DISCUSSION AND LIMITATIONS

The present operator is an estimate-level projection layer. It is graph-local in the sense that each sensor uses only its communication neighborhood, but it is not yet a fully recursive message update inside the LMB filtering loop. This distinction matters for the final paper claim: the current evidence supports label-aligned output fusion under unreliable communication, while an online distributed LMB recursion would require additional stability and lifecycle analysis.

The method also uses Euclidean position matching and equally weighted moment barycenters. These

TABLE I

Fifty-trial paired validation under tiered packet loss. Lower values are better. Values are trial means. The proposed operator improves both network agreement and local tracking quality; the reference-only ablation isolates the part of the gain that comes from label-set canonicalization alone.

Method	Net OSPA	Loc. disag.	Card. disp.	E-OSPA	RMSE	CardErr
Tuned spatial-KLA AA	1.682607	1.472837	0.035350	2.029641	3.682880	0.088000
Neighborhood reference-only	1.030062	0.891585	0.021200	1.911286	3.662955	0.077200
Neighborhood label-barycenter	0.309818	0.216372	0.021200	1.681483	3.449035	0.077200

TABLE II

Paired reductions relative to the tuned spatial-KLA AA baseline. The bracketed interval is the 95% confidence interval of the absolute paired reduction.

Metric	Full reduction	Full wins	Ref.-only reduction	Ref.-only wins
Network OSPA	81.59% [1.359, 1.386]	50/50	38.78% [0.639, 0.666]	50/50
Loc. disagreement	85.31% [1.238, 1.275]	50/50	39.46% [0.566, 0.596]	50/50
Card. dispersion	40.03% [0.012, 0.016]	50/50	40.03% [0.012, 0.016]	50/50
E-OSPA	17.15% [0.342, 0.355]	50/50	5.83% [0.112, 0.125]	50/50
RMSE	6.35% [0.201, 0.267]	48/50	0.54% [-0.017, 0.056]	27/50
Card. error	12.27% [0.009, 0.013]	46/50	12.27% [0.009, 0.013]	46/50

TABLE III

Filter runtime over 50 paired trials.

Method	Runtime (s)	Relative
Tuned spatial-KLA AA	62.245 $\pm$ 6.444	1.000x
Neighborhood reference-only	95.576 $\pm$ 11.190	1.542x
Neighborhood label-barycenter	102.347 $\pm$ 11.262	1.647x

choices make the operator simple and interpretable, but they do not use covariance quality, label age, or link reliability inside the barycenter itself. A covariance-aware or reliability-weighted barycenter may improve robustness when local uncertainty varies strongly across sensors. Finally, the current validation is concentrated on a tiered packet-loss formation scenario. A submission-ready evaluation should add independent verification and at least one broader scenario family.

IX. CONCLUSION

This paper reframes a persistent AA-LMB fusion error as a label-space alignment and matched-posterior barycenter problem. The proposed neighborhood label-barycenter operator chooses a robust local reference, canonically matches neighboring tracks to that reference, and computes first-two-moment barycenters for matched states. In 50 paired trials under tiered packet loss, the method improves both network agreement and local tracking metrics, while a reference-only ablation shows that the main spatial

gain comes from posterior barycentering rather than label copying alone. The next step is to move the operator from an estimate-level projection into a recursive distributed LMB update with explicit lifecycle guards.

REFERENCES

- [1] B.-N. Vo, B.-T. Vo, and D. Phung
Labeled random finite sets and the Bayes multi-target tracking filter
IEEE Transactions on Signal Processing, vol. 62, no. 24, pp. 6554–6567, 2014.
- [2] S. Reuter, B.-T. Vo, B.-N. Vo, and K. Dietmayer
The labeled multi-bernoulli filter
IEEE Transactions on Signal Processing, vol. 62, no. 12, pp. 3246–3260, 2014.
- [3] G. Battistelli and L. Chisci
Kullback-leibler average, consensus on probability densities, and distributed state estimation with guaranteed stability
Automatica, vol. 50, no. 3, pp. 707–718, 2014.
- [4] B. Wang, W. Yi, R. Hoseinnezhad, S. Li, L. Kong, and X. Yang
Distributed fusion with multi-bernoulli filter based on generalized covariance intersection
IEEE Transactions on Signal Processing, vol. 65, no. 1, pp. 242–255, 2017.
- [5] S. Li, W. Yi, R. Hoseinnezhad, G. Battistelli, B. Wang, and L. Kong
Robust distributed fusion with labeled random finite sets
IEEE Transactions on Signal Processing, vol. 66, no. 2, pp. 278–293, 2018.
- [6] T. Li, R. Yan, K. Da, and H. Fan
Arithmetic average density fusion—part III: Heterogeneous unlabeled and labeled RFS filter fusion
IEEE Transactions on Aerospace and Electronic Systems, vol. 60, no. 1, pp. 1023–1034, 2024.
- [7] H. W. Kuhn
The hungarian method for the assignment problem

- Naval Research Logistics Quarterly, vol. 2, no. 1–2, pp. 83–97, 1955.
- [8] K. Shen, P. Dong, Z. Jing, and H. Leung
Consensus-based labeled multi-bernoulli filter for multi-target tracking in distributed sensor network
IEEE Transactions on Cybernetics, vol. 52, no. 12, pp. 12 722–12 733, 2022.
 - [9] L. Gao, G. Battistelli, and L. Chisci
Resilient labeled multi-bernoulli fusion with peer-to-peer sensor network
Information Fusion, vol. 100, p. 101965, 2023.
 - [10] T. Li, Y. Song, E. Song, and H. Fan
Arithmetic average density fusion–part I: Some statistic and information-theoretic results
Information Fusion, vol. 104, p. 102199, 2024.
 - [11] T. Li
Arithmetic average density fusion–part II: Unified derivation for unlabeled and labeled RFS fusion
IEEE Transactions on Aerospace and Electronic Systems, vol. 60, no. 3, pp. 3255–3268, 2024.
 - [12] T. Li, H. Liang, G. Li, J. Garcia Herrero, and Q. Pan
Arithmetic average density fusion–part IV: Distributed heterogeneous fusion of RFS and LRFS filters via variational approximation
IEEE Transactions on Signal Processing, vol. 73, pp. 1454–1469, 2025.
 - [13] T. Li, H. Li, G. Li, and H. Miao
Arithmetic average density fusion–part V: Fisher information weighted fusion of multi-rate filters
TechRxiv, 2026, preprint.
 - [14] L. Zheng, L. Shi, F. Yang, and Y. Cai
Distributed weighted fusion of RFS densities for peer-to-peer sensor networks
Information Fusion, vol. 129, p. 104060, 2026.
 - [15] W. Yi and L. Chai
Heterogeneous multi-sensor fusion with random finite set multi-object densities
IEEE Transactions on Signal Processing, vol. 69, pp. 3399–3414, 2021.
 - [16] L. Gao, G. Battistelli, and L. Chisci
Fusion of labeled RFS densities with different fields of view
IEEE Transactions on Aerospace and Electronic Systems, vol. 58, no. 6, pp. 5908–5924, 2022.
 - [17] C. Cao and Y. Zhao
Distributed heterogeneous multi-sensor fusion based on labeled random finite set for multi-target state estimation using information geometry
Signal Processing, vol. 238, p. 110135, 2026.
 - [18] M. Üney, J. Houssineau, E. Delande, S. J. Julier, and D. E. Clark
Fusion of finite-set distributions: Pointwise consistency and global cardinality
IEEE Transactions on Aerospace and Electronic Systems, vol. 55, no. 6, pp. 2759–2773, 2019.
 - [19] D. Schuhmacher, B.-T. Vo, and B.-N. Vo
A consistent metric for performance evaluation of multi-object filters
IEEE Transactions on Signal Processing, vol. 56, no. 8, pp. 3447–3457, 2008.
 - [20] L. Xiao, S. Boyd, and S. Lall
A scheme for robust distributed sensor fusion based on average consensus
In Proceedings of the 4th International Symposium on Information Processing in Sensor Networks, 2005, pp. 63–70.