

Guarded Current-Evidence Fusion for Cooperative Multitarget Tracking

Abstract—Cooperative posteriors mix inherited and current information, which age weighting cannot control separately. We propose guarded current-evidence fusion (GCE), admitting bounded Bernoulli posterior-to-prediction ratios using sensing opportunity, score and miss support, and Gaussian curvature. Admitted ratios jointly update spatial density and existence through a shared normalizer. On 25 previously used V2V4Real sequences with a common local tracker, GCE reduces OSPA by 6.7%/5.7% over qualified geometric pooling under reliable/intermittent links. Recursive controls show a smaller advantage over equally guarded scalar correction; the correction-guard interaction intervals include zero. GCE has lower mean OSPA on three newly acquired segments, but underperforms pooling on the sole previously absent recording. Correlated Gaussian observations also yield undercovered spatial regions despite positive semidefinite increments. Exact zero omission reduces fragmented traffic by 11.3%/10.8% relative to full GCE packets, with remaining overhead over pooling. These results support selective admission on the evaluated corpus while limiting claims about generalization and uncertainty.

I. INTRODUCTION

Cooperative robots observe different targets, occlusions, and missed detections. Their exchanged tracks also contain information inherited from earlier communication. A recently received posterior may therefore carry old evidence, while a low existence probability may convey a useful current miss. Effective fusion must distinguish inherited beliefs from new observations without transmitting the entire sensing history.

Labeled multi-Bernoulli (LMB) filtering represents each track by a label, existence probability, and conditional spatial density [1], [2]. Geometric pooling, or Kullback–Leibler averaging (KLA), combines local posteriors when their cross-correlations are unavailable [3]. In the ideal shared-prior case, normalized geometric weights retain one copy of the prior but also temper each new likelihood. Changing posterior weights according to observation age does not separately control these contributions. A further coupling arises within each Bernoulli: the integral of its spatial product affects the probability that the target exists. An existence correction alone leaves this part of the current evidence unused. Figure 1 illustrates the distinction: a posterior weight acts on inherited and current information together, whereas a current ratio can change their relative contributions.

We propose *guarded current-evidence fusion* (GCE). Each robot encodes the ratio between its local update and prediction. Fusion augments a conservative posterior pool with admitted fractions of these ratios, then jointly normalizes existence and spatial density. Current sensing opportunity, association-weighted score support, and Gaussian curvature determine which increments are admitted. Figure 2 shows




Fig. 1. Cooperative tracking with shared history (illustration). In the idealized shared-prior case, equal-weight geometric pooling retains the prior but tempers each current likelihood. GCE admits bounded fractions of current posterior-to-prediction ratios, increasing the likelihood exponents ω_j while retaining one prior. Joint normalization carries this correction into both spatial density and existence.

how the same current evidence reaches both branches of the fused belief.

Prior-likelihood separation is established [4], [5], as is normalization of a Bernoulli spatial product. Our focus is deciding which approximate local increments to admit when sensing histories and associations differ. The paper contributes:

- A bounded admission rule combining current opportunity, association-weighted score and miss support, and a source-level curvature check for approximate Gaussian Bernoulli ratios.
- A closed-form implementation that couples admitted spatial and existence corrections, with an exact omission code for zero vectors.
- Recursive correction-curvature controls, a fixed-strength ratio baseline, and evaluations of new sequence inputs, known ego motion, detection-probability sensitivity, and correlated Gaussian observations.

II. RELATED WORK

A. Cooperative robotic tracking

Distributed probability hypothesis density (PHD) filters support robot search [6], decentralized action selection [7],




Fig. 2. Guarded current-evidence fusion. Robots exchange inherited beliefs and current Bernoulli ratios. Qualified support forms the base; sensing, score, and curvature determine the admitted correction. The same ratio updates spatial density and existence through the shared integral I . The fused belief becomes the next local prior. Robot poses and density shapes illustrate the mechanism.

and sensing reconfiguration [8]; LMB filters also support distributed mapping [9]. V2V4Real provides real cooperative driving observations [10], and DMSTrack learns uncertainty in a differentiable multisensor tracker [11].

Where2comm selects spatial features by confidence [12]; CoopTrack exchanges instance features and learns cross-agent association [13]. GCE operates on probabilistic track summaries using DMSTrack’s per-vehicle detections.

B. Finite-set fusion and sensing support

Covariance intersection addresses fusion with unknown cross-correlations [14], and KLA formulates geometric pooling over probability densities [3]. Generalized covariance intersection (GCI) extends geometric pooling to PHD filters [15]; consensus labeled random finite set (RFS) filters preserve track identities [16].

Direct GCI fusion can degrade under inconsistent labels [17]; linear assignment can align local identities before labelwise fusion [18]. Different fields of view require explicit treatment of absent labels and exclusive information [19]–[21]. Minimum-information-loss (MIL) fusion uses arithmetic pooling of labeled finite-set densities [22]; track consensus matches and combines short trajectories [23]. These methods motivate our use of aligned labels and explicit sensing support.

Cardinality and localization also play distinct roles in finite-set fusion: pointwise density consistency alone does not ensure a useful cardinality distribution [24]. Heterogeneous fusion assigns confidence at the component level [25]. GCE instead couples the two components through the normalizer of an admitted Bernoulli update.

C. Prior and current likelihood information

Shared-prior Bayesian fusion corrects for common information in local posteriors [4]. Likelihood consensus aggregates exponential-family likelihood representations [5]. GCE transmits Gaussian natural-parameter increments and admits them fractionally. Local LMB association and moment projection can make these ratios differ from exact sensor likelihoods. The admission rule is a practical restriction on such approximations; it does not estimate unknown cross-source correlation or certify covariance consistency.

III. GUARDED CURRENT-EVIDENCE FUSION

A. Local prediction, update, and information support

Each robot predicts its LMB, updates it with the current local measurement set, and exchanges one posterior packet with available neighbors. Fusion output becomes the next local prior. A label is aligned across sources before fusion; births remain local and measurement driven. For a represented label, write the local prediction and update as $f_j^- = (r_j^-, p_j^-)$ and $f_j^+ = (r_j^+, p_j^+)$. The Bernoulli density has empty-set mass $1-r$ and singleton density $rp(x)$. The implementation moment-projects p_j^- and p_j^+ to four-dimensional Gaussians over planar position and velocity.

A direct sensing opportunity occurs when the modeled detection probability at a predicted track is positive. Both a detection and a missed detection advance the local direct-opportunity timestamp τ_j ; receiving another robot’s packet does not advance it. A separate history flag records whether a represented belief descends from an informative local update. When informed inputs exist, untouched priors are excluded. An absent label contributes an existence censor only if its source could observe the represented target. Otherwise

that source abstains. Censors carry no spatial density. This support handling is shared with the qualified no-age and recency controls.

Let $\mathcal{P}$ be eligible represented sources and $\mathcal{E}$ eligible existence sources, including qualified censors. Normalize ordinary delivered weights w_j separately on the two supports:

$$\alpha_j = \frac{w_j \mathbf{1}[j \in \mathcal{P}]}{\sum_{k \in \mathcal{P}} w_k}, \quad b_j = \frac{w_j \mathbf{1}[j \in \mathcal{E}]}{\sum_{k \in \mathcal{E}} w_k}. \quad (1)$$

All sums below use the corresponding support. A label without spatial support is omitted. For a represented source, define

$$a_j = \rho + (1 - \rho)e^{-(t - \tau_j)\Delta t / \lambda}, \quad q_j = \frac{b_j a_j}{\sum_k b_k a_k}, \quad (2)$$

with $a_j = \rho$ before any local opportunity and $a_j = 1$ for a censor.

B. A conservative inherited base

Set $z_j = \text{logit}(r_j^+)$ and choose the existence weights

$$\beta = \begin{cases} q, & \sum_j (q_j - b_j) z_j < 0, \\ b, & \text{otherwise.} \end{cases} \quad (3)$$

Thus history can lower the base existence log-odds but cannot raise them. The spatial base retains the ordinary weights α . An unnormalized base has masses

$$A_0 = \prod_j (1 - r_j^+)^{\beta_j}, \quad (4)$$

$$A_1(x) = \prod_j (r_j^+)^{\beta_j} \prod_j p_j^+(x)^{\alpha_j}. \quad (5)$$

These are separate existence and spatial weights; when the weights and supports coincide, the base reduces to ordinary Bernoulli KLA. The no-history ablation sets $\beta = b$.

C. Admitting current positive and negative evidence

The local existence increment is

$$\delta_j = \text{logit}(r_j^+) - \text{logit}(r_j^-). \quad (6)$$

The local tracker incorporates a detection score s through a fitted positive-to-negative mark ratio $L(s) > 0$. Belief propagation approximates the marginal association probabilities [26]. Let W_{jm} be the normalized local association mass for assigning this Bernoulli to measurement m , and let $A_j = \sum_m W_{jm}$. For a current sensing opportunity, use

$$g_j^+ = \sum_m W_{jm} \max\left(0, \frac{L(s_m) - 1}{L(s_m) + 1}\right), \quad (7)$$

$$g_j^- = (1 - A_j) \frac{p_{D,j}}{2 - p_{D,j}}. \quad (8)$$

Both gates are zero without a current opportunity. The first admits positive increments in proportion to association-weighted mark discrimination; the second weights the missed-detection branch by its bounded likelihood contrast. They satisfy $g_j^+ \leq A_j$, $g_j^- \leq 1 - A_j$, and hence $g_j^+ + g_j^- \leq 1$.

If at least two eligible sources represent the label and none of the eligible existence inputs is a censor, define

$$\kappa_j = (1 - \beta_j) \begin{cases} g_j^+, & \delta_j \geq 0, \\ g_j^-, & \delta_j < 0. \end{cases} \quad (9)$$

Otherwise set every $\kappa_j = 0$. This restriction avoids constructing an unavailable current ratio for an absent label. The no-score ablation replaces the contrast in (7) by one, retaining the same score-aware local tracker and setting $g_j^+ = A_j$.

D. Joint existence and spatial normalization

For each admitted source, multiply the same fractional Bernoulli ratio $(f_j^+ / f_j^-)^{\kappa_j}$ into both hypotheses in (5). The resulting spatial product is

$$h(x) = \prod_j p_j^+(x)^{\alpha_j} \prod_j \left[\frac{p_j^+(x)}{p_j^-(x)} \right]^{\kappa_j}, \quad I = \int h(x) dx. \quad (10)$$

Bernoulli normalization identity. For $0 < r_j^-, r_j^+ < 1$ and $0 < I < \infty$, normalizing these empty and singleton masses yields

$$p^*(x) = h(x) / I, \quad (11)$$

$$\text{logit}(r^*) = \sum_j \beta_j z_j + \sum_j \kappa_j \delta_j + \log I. \quad (12)$$

Proof. The empty mass is $H_0 = A_0 \prod_j [(1 - r_j^+) / (1 - r_j^-)]^{\kappa_j}$. The integrated singleton mass is $H_1 = I \prod_j (r_j^+)^{\beta_j} \prod_j (r_j^+ / r_j^-)^{\kappa_j}$. Taking $\log(H_1 / H_0)$ gives (12); normalizing the singleton shape gives (11). $\square$

This normalization identity provides no error ordering or covariance-consistency guarantee. Scalar correction keeps the original spatial pool and integral in (12); GCE recomputes both from the admitted ratios.

E. Gaussian correction and curvature guard

For $p = \mathcal{N}(m, P)$, write

$$\log p(x) = -\frac{1}{2} x^\top J x + v^\top x + c, \quad J = P^{-1}, \quad v = Jm, \quad (13)$$

where $c = -[d \log(2\pi) + \log |P| + m^\top Jm] / 2$ and $d = 4$. A source transmits the differences $(\Delta J_j, \Delta v_j, \Delta c_j)$ between its update and prediction. The four-dimensional Gaussian ratio requires fifteen numbers: ten for the symmetric precision, four for the information vector, and one constant.

Moment projection can widen a covariance after ambiguous association, so this Gaussian ratio need not be log-concave. GCE admits a source correction only when its precision increment is positive semidefinite, up to numerical tolerance:

$$\bar{\kappa}_j = \kappa_j \mathbf{1}[\Delta J_j \succeq 0]. \quad (14)$$

The guard acts on the entire ratio, including the scalar

increment. The aggregate natural parameters are

$$J = \sum_j \alpha_j J_j^+ + \sum_j \bar{\kappa}_j \Delta J_j, \quad (15)$$

$$v = \sum_j \alpha_j v_j^+ + \sum_j \bar{\kappa}_j \Delta v_j, \quad (16)$$

$$c = \sum_j \alpha_j c_j^+ + \sum_j \bar{\kappa}_j \Delta c_j. \quad (17)$$

For positive definite J , the normalized quantities are

$$P^* = J^{-1}, \quad m^* = J^{-1}v, \quad (18)$$

$$\log I = c + \frac{1}{2}[d \log(2\pi) - \log |J| + v^\top J^{-1}v]. \quad (19)$$

Use $\bar{\kappa}$ in (12) as well. A nonfinite or ill-conditioned aggregate, or a failed Cholesky factorization, restores the base. In exact arithmetic, positive semidefinite admitted precision gains and a positive definite base ensure a positive definite aggregate. Positive semidefiniteness does not establish that an increment is statistically reliable: correlated measurement errors can pass the same test. The no-curvature ablation removes the source test while retaining the aggregate integrability check. Current ratios and gates are recomputed after each local update. For state dimension d and M eligible sources, the dense Gaussian arithmetic costs $O(Md^3)$ and adds $O(d^2)$ numbers per track.

Current-likelihood interpretation. Suppose all sources share a prior f^- , the supports and weights coincide ($\alpha = \beta$), and each exact local update satisfies $f_j^+(X) \propto f^-(X)\ell_j(X)$ for a Bernoulli set X . Substitution into the joint product gives

$$f^*(X) \propto f^-(X) \prod_j \ell_j(X)^{\omega_j}, \quad \omega_j = \beta_j + \bar{\kappa}_j \in [\beta_j, 1]. \quad (20)$$

The base retains one prior because $\sum_j \beta_j = 1$; each ratio cancels its own prior. Admission therefore restores a bounded fraction of current likelihood without increasing the exponent of the inherited prior. Zero admission gives the original pool. In the formal unit-admission limit, independent observations recover the shared-prior Bayesian combination [4]. With unequal inherited beliefs, GCE uses the directly available local ratios rather than assuming a common prior can be reconstructed. This distinction motivates separating the history switch from the current-evidence gates.

F. Exact zero-vector transmission

The full Gaussian implementation extends a 232-byte scalar Bernoulli summary by fifteen 64-bit numbers, giving 352 bytes per track plus a 32-byte packet header. Many local updates leave all fifteen differences exactly zero. A two-byte tag distinguishes a full vector from an omitted zero vector and preserves the sign bit of each omitted zero. All nonzero vectors retain their original 120 bytes. For n tracks and k nonzero vectors, the actual packet length is

$$B = 32 + 234n + 120k \text{ bytes.} \quad (21)$$

This codec has no threshold or quantization parameter. Decoded numerical values, and hence the fusion rule, are unchanged.

IV. EXPERIMENTS

A. Previously used and newly acquired data

We use DMSTrack's released per-vehicle detections [11] on V2V4Real [10]. Nine test-backed validation sequences form the development subset (1,993 paired frames); 25 train sequences form the larger comparison (5,601 frames). Both groups and earlier train probes informed fusion design, and the detector was trained on train. These are previously used data.

We fixed GCE settings, the detector checkpoint, and calibration before acquiring the official validation archive. It yields three paired segments (748 frames) from two recordings. Two segments share an earlier recording; only the 409-frame segment is from a previously absent one. No scene or sensor file is identical to an earlier one. The checkpoint configuration assigns train/test to training/validation; its complete history is unverified. Replay uses the same checkpoint with CPU voxelization and MPS inference. Twelve old sensor frames reproduce detection counts, with maximum matched center/score differences of 0.0078 m/0.0031. New labels serve geometry and scoring only.

B. Shared local tracking and score likelihood

We track planar car centers within 40 m of either sensor and the annotated ego-frame rectangle $[-70.4, 70.4] \times [-40, 40]$ m, excluding a 3 m disk around each vehicle. Truth and estimates use the same domain; each local sensor uses its own 40 m support. Detections retain the released confidence cutoff of 0.20.

All methods share the score-aware LBP-LMB local update. The primary constant-velocity model uses current ego coordinates, a 0.1 s interval, survival probability 0.99, acceleration intensity $25 \text{ m}^2/\text{s}^3$, position covariance $I_2 \text{ m}^2$, and in-domain $p_D = 0.9$. Poisson clutter has mean three per sensor scan and density $3/(\pi 40^2)$. Each previous-scan detection creates an independent local birth with existence 0.01, zero velocity, and covariance $\text{diag}(16, 16, 225, 225)$. We prune $r \leq 0.001$, moment-project retained mixtures to one Gaussian, and extract MAP cardinality. Geometric methods share history qualification and absent-label censoring, with in-support sensor existence 0.001.

Following post-hoc probability calibration [27], we fit $u(s) = \sigma(a \logit s + b)$, $a \geq 0$, by sequence-balanced binary log loss with regularization $0.001a^2$. Offline per-sensor one-to-one assignment within 12 m supplies labels. Development sequences use leave-one-sequence-out fits; the 25-sequence corpus and all new data use the full-nine fit, $a = 2.8858$, $b = 2.1564$. The score enters each local association likelihood once through

$$L(s) = \frac{u(s)}{1 - u(s)} \frac{1 - \pi}{\pi}, \quad (22)$$

where π is the sequence-balanced fitting prevalence.

C. Communication and baselines

Reliable links deliver both directed packets each scan. Intermittent links combine independent directed loss of 0.1

TABLE I

COMPARISON ON 25 PREVIOUSLY USED SEQUENCES WITH A SHARED LOCAL TRACKER. OSPA IS IN M; MISSED/FALSE GOSPA COSTS ARE IN M². LOWEST VALUES ARE BOLD.

Method	Reliable links			Intermittent links		
	OSPA	Missed	False	OSPA	Missed	False
Local	5.914	190.030	45.556	5.914	190.030	45.556
MIL-AM	4.014	74.566	56.973	4.163	82.428	58.363
TC-5	4.253	68.620	75.289	4.693	101.053	64.733
TC-10	4.414	88.338	65.001	4.829	116.571	57.691
No-age KLA	3.702	85.564	40.715	3.988	92.977	44.776
Recency	3.672	78.330	43.986	3.952	87.113	47.603
Conservative	3.672	86.046	37.436	3.956	93.865	42.483
Capped-score	3.666	83.065	39.074	3.913	90.189	43.663
Capped-cal.	3.650	80.776	40.418	3.897	88.391	44.318
Scalar	3.526	78.625	34.536	3.832	86.342	40.523
Guarded Scalar	3.508	75.839	35.438	3.793	83.328	41.019
Fixed Ratio (0.25)	3.691	77.557	44.368	3.926	83.138	47.040
GCE	3.456	71.095	36.303	3.761	80.344	41.971

TABLE II

RECURSIVE CONTROLS ON THE SAME 25 SEQUENCES. SCALAR, GUARDED SCALAR, W/O CURVATURE, AND GCE FORM THE SCALAR/JOINT $\times$ GUARD-OFF/ON COMPARISON. OTHER ROWS REMOVE ONE ADMISSION COMPONENT; UNITS FOLLOW TABLE I.

Method	Reliable links			Intermittent links		
	OSPA	Missed	False	OSPA	Missed	False
Scalar reference	3.526	78.625	34.536	3.832	86.342	40.523
Guarded Scalar	3.508	75.839	35.438	3.793	83.328	41.019
w/o curvature guard	3.491	74.845	35.376	3.805	84.047	41.615
w/o history switch	3.467	70.282	38.148	3.776	80.266	43.105
w/o score constraint	3.713	55.993	59.014	3.922	69.042	57.784
GCE (complete)	3.456	71.095	36.303	3.761	80.344	41.971

with a bidirectional outage over the middle 20% of scans. Methods share the delivery trace and one packet opportunity per direction per scan. Receivers pool local and delivered neighbor updates with weights $[0.5, 0.5]$; otherwise they keep the local update. Accounting uses 16,384-byte fragments and a 256-byte control charge per scan. Delivery is independent of packet length. GCE uses $\rho = 0.25$ and $\lambda = 5$ s.

No-age KLA uses qualified weights α, b . *Recency* replaces b with q in existence, and *Conservative* uses β without a current correction. *Capped-score/Capped-cal.* bound recency existence by $\min(r_{\text{rec}}, \max(r_{\text{noage}}, U))$, where U is the largest association-weighted raw/calibrated score among currently observed, upweighted sources with existence at least 0.5.

Scalar retains the original spatial pool p_0 and integral I_0 while applying GCE's scalar gates without a curvature test. *Guarded Scalar* (GS) additionally uses the identical source curvature rejection and aggregate fallback as GCE on its own inputs, but returns p_0 and $\logit r = \sum_j \beta_j z_j + \sum_j \bar{\kappa}_j \delta_j + \log I_0$. GS runs its own full recursion. *Fixed Ratio* retains GCE's joint update, guard conditions, and history weights, replacing both sign-dependent supports by η . The lowest two-condition mean OSPA on the nine development

sequences selects $\eta = 0.25$ from $\{0.25, 0.5, 1\}$, then holds it fixed. GS and Fixed Ratio require full Gaussian ratio packets.

MIL-AM adapts the public manuscript's augmented assignment and common/exclusive subspace fusion [28]. Symmetric KL cost with unmatched cost 50 per side aligns labels; geometric methods share this alignment. Common labels use arithmetic existence pooling and existence-weighted spatial mixtures; exclusive labels retain their source density. *TC-5/TC-10* use the author's kinematic matching and two-stage multi-neighbor functions [23], with five-/ten-scan histories, order one, and cutoff 12 m. They retain separate local filters. *Local* omits communication.

D. Metrics and recursive controls

We use $p = 2, c = 12$ m for position OSPA [29] and missed/false squared GOSPA costs [30] ($\alpha = 2$). Common-target RMSE pools robot-time-truth matches shared by both methods. Sequences receive equal weight. Descriptive 95% intervals use 10,000 paired whole-sequence percentile resamples of the old 9/25 cohorts, without multiplicity adjustment. New-segment means are reported without an inferential interval.

Scalar, GS, joint correction without the guard, and GCE form a complete recursive correction-curvature 2×2 compar-




Fig. 3. Model sensitivity. (a,b) Mean OSPA over the same nine development sequences at every modeled detection probability, using all four recursive methods and identical measurements. The two panels share their vertical scale. (c) Coverage of nominal 95% regions in a conditional-existing Gaussian control with 10,000 samples at each correlation. All source curvature tests pass. The dashed line is nominal coverage; unit-admission GCE increasingly undercovers as correlation grows.

ison. We report GCE–GS and the interaction (GCE – GS) – (joint without guard – Scalar). Further ablations remove the history switch or positive score constraint, retaining the shared local score model. A separate fixed-input diagnostic substitutes base/corrected means and integrals at identical GCE-visited inputs; alternate outputs never feed back into its recursion.

E. Motion, detection probability, and correlation

Raw absolute poses match every old sequence and reproduce released relative sensor transforms. After CV prediction and birth creation, known planar yaw/translation maps positions, velocities, and full covariances to the current ego frame, equivalently implementing planar world-frame CV. No-age, Scalar, GS, and GCE run on all 34 old and three new segments. Separately, modeled $p_D \in \{0.7, 0.8, 0.9, 0.95\}$ varies on the nine development sequences. Both checks retain the same measurements, domains, calibration, and fusion settings, without primary reselection.

To distinguish integrability from statistical reliability, we also simulate a two-dimensional, conditional-existing Gaussian state with prior covariance $25I_2$ and two measurements of noise covariance I_2 . Their cross-covariance is ϱI_2 , for $\varrho \in \{0, 0.25, 0.5, 0.75, 0.9\}$, with 10,000 samples per setting. We compare equal posterior pooling, unit-admission GCE assuming independent errors, and the oracle with known joint covariance, reporting RMSE, normalized squared error, and nominal 95% region coverage. This control does not estimate the real tracker’s covariance coverage.

V. RESULTS

A. Fusion on the previously used corpus

GCE has the lowest mean OSPA among the main comparisons (Table I), improving on No-age KLA by 6.7%/5.7% and on Recency by 5.9%/4.8% (reliable/intermittent).




Fig. 4. Communication and tracking accuracy on the 25-sequence corpus. Circles/diamonds denote reliable/intermittent links; open markers are full GCE packets. Arrows show exact zero omission reducing raw payload at identical OSPA. Encoded GCE still costs more than No-age KLA. Lower and leftward are better.

Against No-age, missed-target cost falls by 16.9%/13.6% and false-target cost by 10.8%/6.3%. Common-target RMSE decreases from 0.964 to 0.907 m and from 1.085 to 1.050 m. These are fusion comparisons on the development corpus, not an independent estimate of end-to-end generalization.

GCE–Fixed Ratio is -0.236/-0.164 m. Both paired intervals lie below zero. This supports adaptive score/miss admission relative to the selected constant-strength control; it does not rule out a better fixed strength outside the three-point development grid.

B. What the recursive controls identify

Guarded Scalar narrows the gap previously attributed to joint correction (Table II). GCE–GS is -0.053/-0.032 m, with intervals $[-0.10424, -0.00522]$ and $[-0.06517, -0.00006]$.

TABLE III
MEAN COMMUNICATION PER SEQUENCE (MiB). FRAG. INCLUDES
PACKET FRAGMENTATION AND CONTROL TRAFFIC.

Method	Reliable links		Intermittent links	
	Raw	Frag.	Raw	Frag.
No-age KLA	3.89	8.35	3.88	8.29
Scalar	3.73	8.19	3.89	8.32
GCE (full)	5.72	9.79	5.98	9.95
GCE (encoded)	4.56	8.69	4.72	8.88

TABLE IV
OSPA ON NEWLY ACQUIRED DATA (M). ALL THREE SEGMENTS TOTAL
748 FRAMES; ONLY THE 409-FRAME SEGMENT IS FROM A RECORDING
ABSENT FROM THE EARLIER ARCHIVES. R/I DENOTE
RELIABLE/INTERMITTENT LINKS.

Method	All 3 segments		New recording	
	R	I	R	I
No-age KLA	4.722	4.848	2.633	3.020
Recency	4.754	4.862	2.663	3.100
Scalar	4.494	4.682	2.726	2.961
Guarded Scalar	4.527	4.619	2.809	3.029
w/o curvature	4.487	4.676	2.730	3.014
Fixed Ratio (0.25)	4.555	4.791	2.766	3.146
GCE	4.460	4.596	2.789	3.055

The intermittent upper bound is within 10^{-4} m of zero. On the nine development sequences, the intermittent GCE-GS interval includes zero.

The correction-curvature interaction is -0.018/-0.005 m, with intervals $[-0.064, 0.026]$ and $[-0.065, 0.048]$. Both include zero, providing no distinct interaction benefit or normalizer-only explanation of the gain. Removing curvature allows large position shifts and activates aggregate fallback twice. Removing positive score support trades fewer misses for more false targets; the history and score-removal total-OSPA intervals include zero.

On identical GCE inputs, changing only the spatial mean at corrected existence reduces immediate OSPA by about 0.003 m. The joint base/base-to-corrected/corrected difference is -0.014/-0.010 m, and its intervals include zero, as do those for changing the integral at corrected means. These substitutions hold admission fixed and never change future predictions. They answer a different question from the fully recursive GS comparison.

C. New segments and the generalization limit

Across all three newly acquired segments, GCE has the lowest mean OSPA among the seven evaluated methods (Table IV). However, two segments come from an already represented recording. On the only previously absent recording, GCE worsens OSPA to 2.789/3.055 m, compared with No-age's 2.633/3.020 m and Scalar's 2.726/2.961 m. The aggregate gain therefore does not establish stable generalization to new recordings. All settings and segments are retained.

D. Motion and detection-probability sensitivity

With planar ego-motion compensation, GCE-GS on the 25 sequences is -0.074/-0.050 m, with paired intervals $[-0.134, -0.023]/[-0.091, -0.016]$ below zero. Absolute error does not improve uniformly. On the new recording, GCE reaches 2.946/3.494 m, compared with No-age's 3.003/3.663 m and Scalar's 2.869/3.479 m. It now beats No-age but remains worse than Scalar. These runs change the coordinate model for every method without retuning admission. Across all four modeled p_D values, GCE retains a lower mean than No-age and GS, but Scalar wins at reliable $p_D = 0.8$ and intermittent $p_D \in \{0.7, 0.95\}$ (Fig. 3a,b). Varying a constant p_D cannot validate visibility under occlusion. All sequence-level sensitivity results, including the nine development sequences and each new segment, remain available in the companion evidence.

E. Integrability does not ensure calibrated uncertainty

Every precision increment in the Gaussian correlation control is positive semidefinite. At zero correlation, unit-admission GCE and the known-correlation oracle agree exactly. At $\rho = 0.9$, GCE's nominal 95% regions cover only 80.55% of samples, versus 95.19% for the oracle (Fig. 3c). Its mean normalized squared error per dimension is 1.847. The guard therefore certifies an admissible Gaussian correction, not reliability under unknown cross-source correlation. Covariance coverage with adaptive gates remains unmeasured.

F. Communication cost

The codec preserves full GCE trajectories and reduces raw payload by 20.3%/21.1%. Fragmentation and control traffic reduce this saving to 11.3%/10.8% (Table III). Relative to native No-age KLA, encoded GCE still requires 17.3%/21.5% more raw payload and 4.1%/7.0% more fragmented bytes. GS and Fixed Ratio use the same full 352-byte per-Bernoulli schema as GCE; native Scalar uses 232 bytes. Packet-length-independent loss emulation does not measure wireless latency, channel capacity, or energy.

VI. DISCUSSION AND CONCLUSION

GCE selectively admits approximate current Gaussian Bernoulli ratios on a conservative pool. Recursive controls support a modest corpus-level gain beyond equally guarded scalar correction and a larger gain over the tested constant strength. They do not isolate a normalizer-only benefit.

The sole previously absent recording exposes a generalization failure despite better averages over all new segments. Additional independent recordings and detector splits are needed. Motion and p_D checks address model dependence but do not supply occlusion-dependent likelihoods; the Gaussian control shows that a PSD guard cannot certify covariance consistency. Zero omission recovers part of the additional communication cost. Larger teams and measured radio links remain necessary before broader deployment claims.

ACKNOWLEDGMENT

OpenAI Codex was used to develop the experimental and analysis code, produce the schematic figures and the result figures and tables from saved outputs, and draft and revise Sections I–VI, the abstract, captions, and mathematical exposition. The introductory scene and mechanism overview were designed with OpenAI image generation and reconstructed as editable vector graphics.

REFERENCES

- [1] B.-T. Vo and B.-N. Vo, “Labeled random finite sets and multi-object conjugate priors,” *IEEE Transactions on Signal Processing*, vol. 61, no. 13, pp. 3460–3475, 2013, doi: 10.1109/TSP.2013.2259822.
- [2] S. Reuter, B.-T. Vo, B.-N. Vo, and K. Dietmayer, “The labeled multi-Bernoulli filter,” *IEEE Transactions on Signal Processing*, vol. 62, no. 12, pp. 3246–3260, 2014, doi: 10.1109/TSP.2014.2323064.
- [3] G. Battistelli and L. Chisci, “Kullback–Leibler average, consensus on probability densities, and distributed state estimation with guaranteed stability,” *Automatica*, vol. 50, no. 3, pp. 707–718, 2014, doi: 10.1016/j.automatica.2013.11.042.
- [4] P. Wu, T. Imbiriba, V. Elvira, and P. Closas, “Bayesian data fusion with shared priors,” arXiv:2212.07311, 2022.
- [5] O. Hlinka, O. Sluciak, F. Hlawatsch, P. M. Djuric, and M. Rupp, “Likelihood consensus and its application to distributed particle filtering,” *IEEE Transactions on Signal Processing*, vol. 60, no. 8, pp. 4334–4349, 2012, doi: 10.1109/TSP.2012.2196697.
- [6] P. M. Dames, “Distributed multi-target search and tracking using the PHD filter,” *Autonomous Robots*, vol. 44, no. 3–4, pp. 673–689, 2020, doi: 10.1007/s10514-019-09840-9.
- [7] A. Banerjee and J. Schneider, “Decentralized multi-agent active search and tracking when targets outnumber agents,” in *2024 IEEE International Conference on Robotics and Automation (ICRA)*, 2024, pp. 7229–7235, doi: 10.1109/ICRA57147.2024.10609977.
- [8] R. K. Ramachandran, N. Fronda, and G. S. Sukhatme, “Resilience in multirobot multitarget tracking with unknown number of targets through reconfiguration,” *IEEE Transactions on Control of Network Systems*, vol. 8, no. 2, pp. 609–620, 2021, doi: 10.1109/TCNS.2021.3059794.
- [9] L. Gao, G. Battistelli, and L. Chisci, “Efficient distributed multi-robot SLAM under missed detections and clutter with labeled multi-Bernoulli map,” *Robotics and Autonomous Systems*, vol. 198, p. 105358, 2026, doi: 10.1016/j.robot.2026.105358.
- [10] R. Xu, X. Xia, J. Li, H. Li, S. Zhang, Z. Tu, Z. Meng, H. Xiang, X. Dong, R. Song, H. Yu, B. Zhou, and J. Ma, “V2V4Real: A real-world large-scale dataset for vehicle-to-vehicle cooperative perception,” in *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 13 712–13 722, doi: 10.1109/cvpr52729.2023.01318.
- [11] H.-K. Chiu, C.-Y. Wang, M.-H. Chen, and S. F. Smith, “Probabilistic 3D multi-object cooperative tracking for autonomous driving via differentiable multi-sensor Kalman filter,” in *2024 IEEE International Conference on Robotics and Automation (ICRA)*, 2024, pp. 18 458–18 464, doi: 10.1109/icra57147.2024.10610487.
- [12] Y. Hu, S. Fang, Z. Lei, Y. Zhong, and S. Chen, “Where2comm: Communication-efficient collaborative perception via spatial confidence maps,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 4874–4886, doi: 10.52202/068431-0352.
- [13] J. Zhong, J. Wang, J. Xu, X. Li, Z. Nie, and H. Yu, “CoopTrack: Exploring end-to-end learning for efficient cooperative sequential perception,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025, pp. 26 954–26 965.
- [14] S. J. Julier and J. K. Uhlmann, “A non-divergent estimation algorithm in the presence of unknown correlations,” in *Proc. American Control Conference*, vol. 4, 1997, pp. 2369–2373, doi: 10.1109/ACC.1997.609105.
- [15] M. Uney, D. E. Clark, and S. J. Julier, “Distributed fusion of PHD filters via exponential mixture densities,” *IEEE Journal of Selected Topics in Signal Processing*, vol. 7, no. 3, pp. 521–531, 2013, doi: 10.1109/JSTSP.2013.2257162.
- [16] C. Fantacci, B.-N. Vo, B.-T. Vo, G. Battistelli, and L. Chisci, “Consensus labeled random finite set filtering for distributed multi-object tracking,” arXiv:1501.01579, 2015.
- [17] S. Li, W. Yi, R. Hoseinnezhad, G. Battistelli, B. Wang, and L. Kong, “Robust distributed fusion with labeled random finite sets,” *IEEE Transactions on Signal Processing*, vol. 66, no. 2, pp. 278–293, 2018, doi: 10.1109/TSP.2017.2760286.
- [18] S. Li, G. Battistelli, L. Chisci, W. Yi, B. Wang, and L. Kong, “Computationally efficient multi-agent multi-object tracking with labeled random finite sets,” *IEEE Transactions on Signal Processing*, vol. 67, no. 1, pp. 260–275, 2019, doi: 10.1109/TSP.2018.2880704.
- [19] —, “Multi-sensor multi-object tracking with different fields-of-view using the LMB filter,” in *2018 21st International Conference on Information Fusion (FUSION)*, 2018, pp. 1201–1208, doi: 10.23919/ICIF.2018.8455250.
- [20] G. Li, G. Battistelli, W. Yi, and L. Kong, “Distributed multi-sensor multi-view fusion based on generalized covariance intersection,” *Signal Processing*, vol. 166, p. 107246, 2020, doi: 10.1016/j.sigpro.2019.107246.
- [21] L. Gao, G. Battistelli, and L. Chisci, “Fusion of labeled RFS densities with different fields of view,” *IEEE Transactions on Aerospace and Electronic Systems*, vol. 58, no. 6, pp. 5908–5924, 2022, doi: 10.1109/TAES.2022.3182642.
- [22] —, “Fusion of labeled RFS densities with minimum information loss,” *IEEE Transactions on Signal Processing*, vol. 68, pp. 5855–5868, 2020, doi: 10.1109/TSP.2020.3028496.
- [23] H. V. Nguyen, H. Rezaatofghi, B.-N. Vo, and D. C. Ranasinghe, “Distributed multi-object tracking under limited field of view sensors,” *IEEE Transactions on Signal Processing*, vol. 69, pp. 5329–5344, 2021, doi: 10.1109/tsp.2021.3103125.
- [24] M. Uney, J. Houssineau, E. Delande, S. J. Julier, and D. E. Clark, “Fusion of finite-set distributions: Pointwise consistency and global cardinality,” *IEEE Transactions on Aerospace and Electronic Systems*, vol. 55, no. 6, pp. 2759–2773, 2019, doi: 10.1109/TAES.2019.2893083.
- [25] W. Yi and L. Chai, “Heterogeneous multi-sensor fusion with random finite set multi-object densities,” *IEEE Transactions on Signal Processing*, vol. 69, pp. 3399–3414, 2021, doi: 10.1109/TSP.2021.3087033.
- [26] J. L. Williams and R. A. Lau, “Approximate evaluation of marginal association probabilities with belief propagation,” *IEEE Transactions on Aerospace and Electronic Systems*, vol. 50, no. 4, pp. 2942–2959, 2014, doi: 10.1109/TAES.2014.120568.
- [27] A. Niculescu-Mizil and R. Caruana, “Predicting good probabilities with supervised learning,” in *Proc. 22nd International Conference on Machine Learning (ICML)*, 2005, pp. 625–632, doi: 10.1145/1102351.1102430.
- [28] L. Gao, G. Battistelli, and L. Chisci, “Fusion of labeled RFS densities with minimum information loss,” arXiv:1911.01083v1, 2019.
- [29] D. Schuhmacher, B.-T. Vo, and B.-N. Vo, “A consistent metric for performance evaluation of multi-object filters,” *IEEE Transactions on Signal Processing*, vol. 56, no. 8, pp. 3447–3457, 2008, doi: 10.1109/TSP.2008.920469.
- [30] A. S. Rahmathullah, A. F. Garcia-Fernandez, and L. Svensson, “Generalized optimal sub-pattern assignment metric,” in *2017 20th International Conference on Information Fusion (Fusion)*, 2017, pp. 1–8, doi: 10.23919/ICIF.2017.8009645.